EMERGENT COORDINATION IN MULTI-AGENT LANGUAGE MODELS

Christoph Riedl

D'Amore-McKim School of Business, Khoury College of Computer Sciences, and Network Science Institute Northeastern University Boston, MA 02118, USA c.riedl@northeastern.edu

ABSTRACT

When are multi-agent LLM systems merely a collection of individual agents versus an integrated collective with higher-order structure? We introduce an information-theoretic framework to test—in a purely data-driven way—whether multi-agent systems show signs of higher-order structure. This information decomposition lets us measure whether dynamical emergence is present in multiagent LLM systems, localize it, and distinguish spurious temporal coupling from performance-relevant cross-agent synergy. We implement both a practical criterion and an emergence capacity criterion operationalized as partial information decomposition of time-delayed mutual information (TDMI). We apply our framework to experiments using a simple guessing game without direct agent communication and only minimal group-level feedback with three randomized interventions. Groups in the control condition exhibit strong temporal synergy but only little coordinated alignment across agents. Assigning a persona to each agent introduces stable identity-linked differentiation. Combining personas with an instruction to "think about what other agents might do" shows identity-linked differentiation and goal-directed complementarity across agents. Taken together, our framework establishes that multi-agent LLM systems can be steered with prompt design from mere aggregates to higher-order collectives. Our results are robust across emergence measures and entropy estimators, and not explained by coordination-free baselines or temporal dynamics alone. Without attributing human-like cognition to the agents, the patterns of interaction we observe mirror well-established principles of collective intelligence in human groups: effective performance requires both alignment on shared objectives and complementary contributions across members.

1 Introduction

Recent advances in generative AI (specifically LLMs) have led to tremendous advances in multiagent systems (Qu et al., 2024; Hong et al., 2024; Qian et al., 2023; Li et al., 2024a; Subramaniam et al., 2025). Multi-agent systems often show impressive performance increases over single-agent solutions (Wu et al., 2023; Chen et al., 2023; Li et al., 2024b; Tao et al., 2024). A key argument behind such performance gains are claims to "greater-than-the-sum-of-its-parts" effects from differentiated agents (Chen et al., 2023, p.1). Connecting multiple differentiated agents has the potential benefit to leverage unique contributions that would not be available if the task was assigned to only a single agent (Fazelpour & De-Arteaga, 2022; Tollefsen et al., 2013; Luppi et al., 2024). Conceptually, any group gains depend on emergence and synergy (Theiner, 2018; Riedl et al., 2021; Page, 2008). Despite impressive performance of many multi-agent systems we do not yet have a principled understanding when and how such synergy emerges, what role agent differentiation plays, and how to steer it systematically. This all points to a crucial need for deeper understanding of collective intelligence in multi-agent systems.

¹This is true for both human groups and multi-agent systems.

In this paper, we address the question when is a population of LLM agents a mere collection vs. a higher-order collective? We follow a purely data-driven approach to assess whether multi-agent systems show signs of emergence in the form of higher-order synergy characterized by structural coupling and joint information about future states and task outcomes. Intuitively, synergy refers to information about a target that a collection of variables provide only jointly but not individually (Rosas et al., 2020; Humphreys, 1997). We develop a principled framework (Figure 1a) building on new insights in information theory (Rosas et al., 2020; Mediano et al., 2022). Our goal is to develop methods to measure emergent role specialization, determine whether multi-agent systems show such signs of emergent synergy, explore whether synergy enables increased performance, and explore whether emergent behavior can be systematically steered with prompting.

We apply our framework to study multi-agent systems of gpt-4.1 and Llama-3.1-8B agents solving a simple group guessing task (Figure 1b) with three treatment interventions: a control condition, a condition that assigns personas to each agent, and one that uses personas with an instruction to "think about what other agents might do" (a theory of mind (ToM) prompt). The paper is organized into addressing three concrete research questions:

- RQ1: Do multi-agent LLM systems possess the capacity for emergence?
- RQ2: What functional advantages—such as synergistic coordination and higher performance—arise when multi-agent systems exhibit emergence?
- RQ3: Can we design prompts, roles, or reasoning structures that steer the internal coordination of multi-agent to encourage positive, goal-directed synergy?

Our findings provide evidence for multi-agent LLM system capacity for emergence and that it underpins performance. Through a variety of complementary analyses and robustness tests, we show that coordination style across interventions is very different. While emergence is present in all, only the ToM-prompt condition leads to groups with identity-linked differentiation and goal-directed complementarity across agents: they operate as an integrated, goal-directed unit. This is consistent with research on human groups showing that performance gains are not automatic (e.g., Stasser & Titus, 1985) and that complementarity needs to be balanced with integration and goal alignment (Theiner, 2018; DeChurch & Mesmer-Magnus, 2010; Riedl et al., 2021).

This paper helps us understand when and how multi-agent systems exhibit higher-order properties, how to control them, and their internal coordination. This can inform multi-agent system design by showing how to combine agents effectively. We make four contributions:

- 1. Novel framework to quantify emergent properties in multi-agent systems based on information decomposition including conditional/residual variants and outcome-relevant partial information decomposition.
- 2. Principled diagnostic approaches to localize where synergy resides and distinguish it from alternative explanations (such as heterogeneous learning rates). Specifically, we develop two surrogate null distribution tests (row-shuffle to probe identity-locked structure; column-shuffle for dynamic alignment). This design lets us tease apart "good" synergy aligned with task goals from spurious or misaligned synergy.
- 3. Demonstrate how to steer emergence with specific prompts. Prompt-level manipulations causally change higher-order dependencies and reliably induces distinct coordination regimes, shifting collectives from spurious and misdirected synergy to stable & goal-aligned complementarity driven by differentiated identities.
- 4. Internal coordination is measurable and controllable with interventions. Groups differ in variance, stability, and adaptability—properties relevant to reliability and deployment.

2 Method

Group Task. We study a group guessing game (developed by Goldstone et al., 2024, who called it group binary search) without communication: agents propose integers whose sum needs to match a randomly generated hidden target number. Agents are unaware of each others' guesses and the size of their group, and receive only group-level feedback "too high" or "too low." The task is

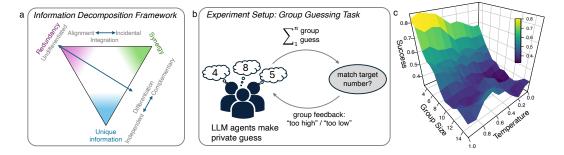


Figure 1: **a)** Information decomposition provides framework to explain tension in multi-agent systems. Agents are either undifferentiated or differentiated, provide independent or complementary information, which is either well aligned or incidental (adapted from Luppi et al., 2024). **b)** Experiment setup of the group binary search task. **c)** Preliminary experiments testing different group sizes and temperature settings. Surface values were smoothed using a local 3×3 weighted averaging filter, giving higher weight to each cell's original value to reduce noise while preserving local structure.

challenging because identical strategies induce oscillation and only complementary strategies yield success. This setting naturally pits redundancy (alignment) against synergy (useful diversity), and admits clear nulls via row-wise (break identities) and column-wise (break alignment) surrogates. Research on humans has shown that groups can reliably solve this task through the emergence of specialized, complementary roles (Goldstone et al., 2024).

Analytical Framework. How would we know if a multi-agent system shows emergent properties that could suggest that the sum is more than its parts? We develop a principled test based on a recent formal theory of dynamic emergence based on information decomposition (Rosas et al., 2020; Mediano et al., 2022; Bedau & Humphreys, 2008). Our framework connects emergence with information about a system's temporal evolution—future states of the whole—with information that cannot be traced to the current state of its parts. This framework is based on partial information decomposition (PID; Williams & Beer, 2010) and time-delayed mutual information (TDMI; Luppi et al., 2022), is data-driven, quantifiable, and amenable to empirical testing with falsifiable conjectures against specific null hypotheses. Combining this with permutation tests, focused either on breaking connections between agents or connections across time allow us not only to quantity emergent properties but also to localize them in the system. This is crucial for multi-agent systems because we care about distinguishing synergy among specialized, differentiated agents from mere dynamic alignment.

We implement three tests. The first, termed the *practical criterion*, concerns predicting a macro signal V of the system. It asks whether the macro contains predictive information beyond any individual part. Given microstate $X_t = (X_{1,t}, \dots, X_{n,t})$ and macro $V_t = f(X_t)$. We align samples $(t,t+\ell)$ and compute the score as

$$S_{\text{macro}}(\ell) = I(V_t; V_{t+\ell}) - \sum_{k=1}^{n} I(X_{k,t}; V_{t+\ell}).$$

A positive value indicates that the macro's self-predictability exceeds what the sum of its parts can explain, thus indicating emergent dynamical synergy. We refer to "dynamical" synergy because targets are time-lagged and do not claim causal directionality beyond the temporal ordering (see Rosas et al., 2020, for a nuanced discussion). It is a coarse, order-agnostic screen sensitive to multipart synergy (i.e., synergy of any order ≥ 2). However, it is also penalized by redundancy across parts which can make the score negative even when higher-order synergy exists.

The second test, termed *emergence capacity*, captures ability of the multi-agent system to host *any* synergy. It measures the predictive synergy from two agents' current states to their *joint* future state. For each pair of agents (i,j) and time t, let the sources be $X_{i,t}$ and $X_{j,t}$, and define the bivariate target as the next-step joint state $T_{ij,t+\ell} \equiv (X_{i,t+\ell}, X_{j,t+\ell})$. From the joint contingency table over $(X_{i,t}, X_{j,t}, T_{ij,t+\ell})$ we compute a two-source PID of the predictive information

$$I(\lbrace X_{i,t}, X_{j,t} \rbrace; T_{ij,t+\ell}) = UI_i + UI_j + Red_{ij} + Syn_{ij}$$

and take Syn_{ij} as the pairwise dynamical synergy. A positive score $\mathrm{Syn}_{ij}>0$ indicates that the macro captures predictive information about the joint future not recoverable from any single component. We compute this for all unordered pairs (i,j) and take the median as group-level synergy capacity. Compared to the practical criterion, this test is limited to detect synergy of order k=2 but has the advantage that it does *not* rely on the definition of a suitable whole-system macro signal.

Finally, we implement a *coalition test*. Let I_3 be the mutual information between a triplet's current state and the future macro signal: $I_3 = I((X_{i,t}, X_{j,t}, X_{k,t}); V_{t+\ell})$. I.e., how much the three agents jointly predict the macro' future. High I_3 indicates that agents' behaviors are coherently organized toward the shared macro, whereas low I_3 signals weak alignment or uncoordinated behavior. Then

$$G_3 = I_3 - \max(I_{2\{1,2\}}, I_{2\{1,3\}}, I_{2\{2,3\}})$$

measures the additional predictive information the full triplet provides over the most predictive pair I_2 (a form of whole-minus-parts metric; Barrett & Seth, 2011; Mediano et al., 2025). If we find $G_3>0$ this means no pair is sufficient to capture information that the triplet contains about the macro signal at $t+\ell$. We again compute I_3 and G_3 for all possible triplets and take the median as group measure. This measure complements the other two by offering a coalition-level test geared towards functional relevance: it helps answers whether the joint information provided by coalitions of agents is actually about the goal. It allows us to localize where macro predictability depends on beyond-pair structure and to rule out "explained by best pair" explanations.

Estimation Details. As microstates $X_t = (X_{1,t}, \ldots, X_{n,t})$ we use each agent *i*'s guess at time t and transform it into the deviation from equal-share contribution to hit the target number: $devs_{i,t} = raw_{i,t} - target/N$. As macro signal V_t we take the group error $\sum raw_{i,t} - target$ (equivalently $\sum devs_{i,t}$). This minimal transformation removes the trivial level differences imposed by the target and aligns agents on a comparable scale, so that remaining variation reflects coordination (who compensates for whom). Over-/under-contributions are interpretable as the sign of deviations.

We use the Williams–Beer two-source PID (Williams & Beer, 2010) with the I_{min} (minimum specific information) redundancy, estimated via plug-in probabilities. For robustness, we also use Jeffreys-smoothed probabilities, and alternatively MMI redundancy (minimum of whole mutual informations; Mediano et al., 2025) with MI estimated with Miller–Madow bias correction (Miller, 1955). All variables are discretized into K=2 bins using quantile binning. Data are encoded as factors with fixed levels 1,2 (and for the joint target as the Cartesian product levels) to avoid dropping empty categories. We report $\ell=1$ which is suitable to detect next-step oscillating behavior (see task description). Because trials end at success, episode lengths vary (see plot of rounds to success in Appendix). We report additional sensitivity tests using "early synergy" truncated to a fixed horizon of $H \in \{10,15\}$ rounds to ensure comparability (results in Appendix A.2 and A.11; main analyses use full trajectories).

Falsification Tests. We perform two different falsification tests of each of the three criteria introduced above. Our primary test assesses significance by comparing the observed value against a null distribution computed on permuted (shuffled) instances of the *devs* data. We use two surrogates: row-wise shuffles (to break identities) and column-wise time-shift surrogates (preserve individual dynamics while disrupting cross-agent alignment). Full details in Appendix. We combine p-values across independently simulated groups (independent seeds, independently sampled targets, and no shared state) using Fisher's method. As a robustness test, we also compute bias-corrected (BC) estimates using time-demeaned data against a block-shuffled null with $\ell=2$ to mitigate autocorrelation confounds. We use both (a) linear regression demeaning and (b) a functional baseline without between-agent synergy (see Appendix for details). These tests provide additional robustness to the $\ell=2$ block size in the shuffled null and auto-correlation concerns.

Entropy Estimation. Small-sample entropy estimation is often challenging because the true distribution is only partially observed and many outcome categories receive few or even zero counts ("empty bins"). Such finite-sample estimation of information measures is biased upward and can yield spurious positive findings. Bias grows with more bins, dimensionality (e.g., using triplets

²We also performed additional sensitivity analyses using alternative specification of the macro signal as first principle component of individual guesses as used in Rosas et al. (2020).

instead of pairs), and small N (few timesteps). The empirical plug-in estimator replace true probabilities with sample frequencies and unseen (or under-sampled) outcomes often are assigned zero or very low probabilities (Hausser & Strimmer, 2009). To account for this issue, we follow best practices and take several steps. First, we compute the emergence capacity and the coalition tests as order k=2 instead of the entire system of n agents. Second, we apply quantile binning (with two bins) to reduce the dimensionality of our data (sensitivity analyses with three bins reported in A.8). Third, we use bias-corrected entropy estimator with Jeffreys' prior, which smooths estimates and avoids empty bins by adding $\alpha=\frac{1}{2}$ pseudo-counts in the Dirichlet estimator (Jeffreys, 1946). This reduces systematic inflation and cures zero-count pathologies. Finally, as a robustness check, we also report results using the Miller–Madow bias-corrected estimator (Miller, 1955), and the MMI redundancy measure (MMI typically overestimates redundancy, making synergy estimates more conservative; Mediano et al., 2025).

Test of Agent Differentiation. We complement the information theoretic measures with a test for agent differentiation based on hierarchical (mixed) models (Gelman & Hill, 2007). We estimate a sequence of three models for each multi-agent experiment:

$$\begin{split} m_0: \quad y_i &= \beta_0 + u_{\mathsf{time}[i]} + \epsilon_i \\ m_1: \quad y_i &= \beta_0 + u_{\mathsf{time}[i]} + u_{\mathsf{agent}[i]} + \epsilon_i \\ m_2: \quad y_i &= \beta_0 + u_{\mathsf{time}[i]} + u_{\mathsf{agent}[i],0} + u_{\mathsf{agent}[i],\mathsf{time}[i]} + \epsilon_i \end{split}$$

where y_i are $devs_{i,t}$, $u_{\text{time}[i]}$ are random intercepts for (continuous) time, $u_{\text{agent}[i]}$ are agent-level random intercepts, and $u_{\text{agent}[i], \text{time}[i]}$ are agent-level random slopes varying by time. The random intercepts capturing differentiation in agents based on their level differences and varying slopes capture heterogeneous learning rates. Using equal-share deviations removes target-level drift and makes agents directly comparable round-by-round. The random intercepts for time captures roundto-round shifts due to group-wide feedback and oscillation. Hence, agent effects reflect relative, identity-linked behavior rather than shared dynamics. The partial pooling of hierarchical models regularizes per-agent estimates, improving stability in short time series. We then use likelihood ratio tests to compare the nested hierarchical models. Comparing $m_0 \to m_1$ asks whether agents differ in their group contribution (some agents guess higher/lower than others), while $m_1 \rightarrow m_2$ tests whether agents vary in how much their contribution changes across rounds (learning rates). A significant p-values (e.g., below conventional 0.05 levels) indicates that the more complex model explains the data significantly better, suggesting the added random effects (e.g., agent-level differences or varying slopes) are meaningful in this group. This offers an interpretable, non-information theoretic test of monotonic learning-rate heterogeneity. Furthermore, this test does not require discretization, giving a complementary failure mode to the information-theoretic analyses.

Together, these four tests of our framework allow us to (a) detect higher-order structure in multiagent systems, (b) assess whether it is driven by redundancy or synergy, (c) localize whether it is identity-locked vs. dynamic alignment, and (d) show whether the higher-order structure is functionally useful for the task. No single measure does all four. Combined with the prompt-level interventions, the framework allow us to test whether interventions (**Plain**, **Persona**, **ToM**) causally increases dynamic synergy, how it affects identity-locked differentiation, and goal alignment (i.e., increase in I_3).

3 EXPERIMENTS

3.1 PRELIMINARY EXPERIMENTS: GROUP SIZE AND TEMPERATURE

To see if LLMs can solve this task and how sensitive results are to different group sizes and temperature settings we ran a first set of preliminary experiments (see Appendix for prompt). We used a frontier-class model, OpenAI's gpt-4.1-2025-04-14. We varied group size from 3 to 15, and temperature from [0,1] in steps of 0.1. For each grid point we ran 50 group experiments (13 group sizes \times 11 temperature settings \times 50 groups = 7,150 experiments; Figure 1c). Multi-agent systems of gpt-4.1 agents can solve the task reliably but encounter substantial challenges. We fit a logistic regression, predicting the probability of success from group size and temperature. We find the task is significantly easier for smaller groups: each additional group member decreased the odds by roughly

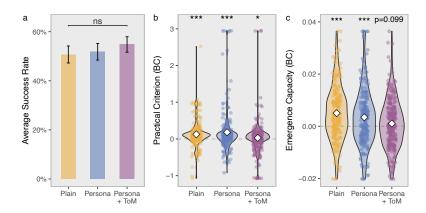


Figure 2: **a)** Group success across three interventions. **b)** Practical emergence criterion (bias corrected). **c)** Emergence capacity dynamical synergy (bias corrected). Data Winsorized at the 1st and 99th percentiles for visual clarity. Stars indicate significance level of Wilcoxon test. Notes: *** p < 0.001; ** p < 0.01; ** p < 0.05.

8% (OR = 0.92, $p < 10^{-16}$). This replicates the same results as for human groups (Goldstone et al., 2024). Each unit increase in temperature increased the odds of success by approximately 50% (OR = 1.50, $p < 10^{-7}$).

3.2 MAIN EXPERIMENTS

For the main set of experiments, we chose groups of N=10 because that appears to be the most difficult setup, and a temperature of 1, and used OpenAI's gpt-4.1-2025-04-14 (OpenAI, 2025) and Meta's Llama-3.1-8B (Meta, 2024) with the prompts shown in Appendix A.1. We query gpt-4.1 through the OpenAI API and Llama through locally hosted Ollama with a NVIDIA RTX A6000. We replicate each group experiment 200 times per treatment condition (600 experiments total). Overall success rate is not significantly different across the interventions (Figure 2a).

Interventions. Plain: The control condition uses only instructions for the guessing game, similar to those used in human subject experiments (Goldstone et al., 2024). Persona: For the persona condition, we follow recent insights into the use of personas within the LLM community (Chen et al., 2024) and ensure personas have relevant attributes: name, gender, age, occupation or skill background, pronouns, personality traits (Goldberg, 1992), and personal values (Cieciuch & Davidov, 2012). **ToM**: The ToM condition instructs agents to think about what other agents might do, and how their actions might affect the group outcome, inspired by chain-of-thought reasoning (Wei et al., 2022). See Appendix A.1 for all prompts, full persona generation recipes, and example persona.

4 RESULTS

4.1 EMERGENT SYNERGY

We first test if there are signs of emergence using the practical emergence criterion (predicting a time-lagged macroscopic group signal from micro-level states). First, we take the criterion computed using the equal-share contribution data and compare it against a null distribution of block-shuffled permutations. Individually, about 3.5% of experiments show a p-value below 0.05. A joint Fisher test using all p-values is highly significant, both individual and when tested within in each treatment condition separately (p-values are below 10^{-16}). Our second test takes the bias-corrected versions of the measure and performs a Wilcoxon signed rank test of the null hypothesis that the median value is above 0 (Figure 2b). We report our most conservative results using time-trend demeaned data. The bias-corrected estimates are well above 0 in all conditions (Plain: $p = 1.5 \times 10^{-16}$; Persona: $p = 6.6 \times 10^{-7}$; ToM: p = 0.02). We report additional robustness tests using (a) raw

data, (b) functional null model residuals, (c) using agent reactivity as data, and (d) bias correction using Jeffrey's smoothing and Miller-Madow in the Appendix.

We repeat analyses using the emergence capacity criterion. We find it is significantly above its null across conditions; results hold under residualization and alternative entropy corrections (Figure 2c; see Appendix for details). Taken together, across two different tests (one using a macro signal the other using future system behavior) and across a variety of robustness tests we find evidence of dynamic emergence capacity in multi-agent LLM systems (answering **RQ1**).

Next, we explore whether agents evolve specialized roles & identities. Are there detectable between-agent differences in either the level of their contribution to the group guess or their temporal evolution (i.e., learning rate)? Using the hierarchical model comparison test we find that agents in may groups have differentiated identities (Figure 3a; a Fisher test for joint *p*-value evidence is again highly significant). There is substantially more differentiation among agents in the **Persona** condition, and even more in the **ToM** condition. Inspecting the reasoning traces of agents, they often contain references to the "personal experience" of their assigned persona such as "In my experience (whether it' corraling cattle or wrangling numbers), starting toward the middle gives the group the best shot to zero in after the first feedback" (see Appendix for additional sample quotes).

4.2 What does Emergence Allow?

Internal Organization of Groups beyond Synergy. What functional differences—if any—does emergence allow? To explore if the internal organization of behavior in groups is different across treatment conditions, we contrast total mutual information (I_3) with information contained in triplets (G_3) . We find significant above zero dynamic emergent synergy in triplets that is not reducible to pairs in all three treatment conditions (again the p-value test shows 20% of groups below 0.05 and a highly significant Fisher test). This provides further evidence for multi-agent LLM system capacity for emergence. G_3 is significantly lower in the ToM condition compared to the Plain condition (two-sample KS test p=0.036) but not different from the Persona condition (p=0.842; Figure 3b). However, we find significantly higher I_3 in the ToM condition compared to both other conditions (both p<0.001; Figure 3c). Finally, we test how many groups have significantly positive I_3 . We find ToM condition has substantially more such groups (Figure 3d). In Appendix A.9 we provide additional robustness tests ruling out heterogeneous learning rates as the sole source of differentiation and provide additional evidence for stable agent identities that persist over time.

Increased Performance. We use regression analysis to explore the joint effect of synergy and redundancy (based on emergence capacity formulation) while carefully controlling for endogeneity of variables measuring emergent properties and performance and controlling for sample selection bias with stabilized inverse probability weighting (see Appendix A.10 for method details). On their own, higher levels of either synergy or redundancy do not predict success. However, when both are present, performance improves significantly (significant interaction with $\beta = 0.24$; p = 0.014). In marginal-effect terms, redundancy amplifies the benefit of synergy on the log-odds scale by 27%; and vice versa synergy amplifies benefits of redundancy by 27%. This pattern implies that systems benefit when redundant pathways create goal alignment, while synergistic interactions extract novel, non-overlapping information—together enabling higher overall performance. We complement the regression analysis with causal mediation analysis (Imai et al., 2010). While reaching only marginal significant levels the effect is consistent: the ToM treatment causally increases performance indirectly by increasing synergy (ACME = 0.034 [95%CI: -0.000-0.07], p = 0.053). This aligns with the interpretation that performance benefits emerge when systems achieve both redundancy (aligned toward a common goal) and synergistic integration (differentiated, complementary roles)—a form of functional complexity (Varley et al., 2023; Luppi et al., 2024).

Together, the results suggest a nuanced picture of emergence in multi-agent systems. Persona and ToM produce agents with distinct, stable identities—moving agents on the differentiated-undifferentiated axis in Figure 1a. This differentiation shapes the space in which other agents can adapt, and specialization by one agent constrains the (useful) actions of the others (agents mutually constrain each other; cf. Tollefsen et al., 2013)—shifts on the interdependent-complementary axis. When combined with ToM, mutual adaptation both strengthens agent differentiation (creating the foundation for possible synergistic complementarity) while also increasing mutual alignment on shared goals (the target signal)—shifts on the integration axis. Differentiation enables synergy,

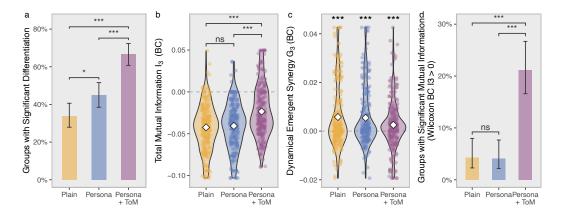


Figure 3: **a)** Agent differentiation using hierarchical mixed model comparison (counting groups in which at least one test (different intercepts or slopes) is below p < 0.05. **b)** Total time-delayed mutual information of triplets (I_3) . **c)** Dynamic emergent synergy (G_3) . **d)** ToM-prompt condition has substantially more groups with significant I_3 content (above 0). In panel a) and d), error bars show Wilson confidence intervals for binary data. Stars indicate significance level of test for equal proportion. Panel b) and c) shows raw data with Jeffreys' correction. Data are Winsorized at the 1st and 99th percentiles for visual clarity. Stars indicate significance level of Wilcoxon test. Notes: *** p < 0.001; ** p < 0.01; * p < 0.05.

but without redundancy and integration, synergy alone does not translate into better collective performance. Agents in the **ToM** condition develop specialized, complementary roles while achieving stronger shared goal alignment: they are simultaneously more differentiated yet also more integrated along highly overlapping task-relevant information which the capacity for synergy can now exploit in productive ways (answering **RQ2** and **RQ3**).

4.3 EXPERIMENTS USING LLAMA-3.1-8B

Repeating the experiments using another class of LLM with 8B parameters (Llama-3.1) we find multi-agent systems are generally not able to solve the task (at least not with the identical prompt as used in the other experiments). Only 10% of groups manage to solve the task (**Plain**: 10.5%; **Persona**: 14%; **ToM**: 5.5%). Model capacity to support reasoning about other agents seems crucial and despite the instruction to do so, Llama-3.1-8B agents fail executing it effectively. Groups in the **ToM** condition perform significantly worse than in the **Plain** condition (test for equality of proportions p = 0.007) thus reversing the findings of the more capable gpt-4.1 agents. Results regarding emergence capacity are mixed (see Appendix).

The strong oscillating pattern in guesses materializes as high predictive power over the group outcome (15.8% groups show significant I_3 compared to only 2% in gpt-4.1 groups) but much lower G_3 synergy between agents (significant synergy present in 2.6% Llama groups compared with 20% of gpt-4.1 groups). There is extremely strong coupling across time (current individual states predict future joint states well as agents oscillate) but very little cross-agent complementarity. This suggests that while multi-agent systems with lower-capacity LLMs still show some signs of emergence, it is mostly spurious across the temporal dimension rather than the productive cross-agent dimension found in gpt-4.1 agents. Lower ToM reasoning reasoning capacity in Llama-3.1-8B (Xiao et al., 2025) may explain the lack of useful, goal-directed cross-agent synergy. Capacity-matched prompt tuning (e.g., lower temperature or simplified instructions) might yield different results.

5 RELATED WORK

Shortly after the release of ChatGPT, the community began exploring multi-agent LLM systems, where multiple LLM agents interact with each other. One early example is Park et al. (2023), who simulated a small population of "individuals" in a *The Sims*-style environment with friends, houses, and jobs, showing the emergence of rich social behaviors such as coordinating a Valentine's Day

party. This was followed by multi-agent systems for complex tasks such as software development (Chen et al., 2023; Hong et al., 2024; Qian et al., 2023), healthcare (Li et al., 2024a), and other domains (Subramaniam et al., 2025). Federated systems or "society of models" also gained attention (Juneja et al., 2024; Li et al., 2024b). These systems often achieve significant performance gains over single-agent baselines (Wu et al., 2023; Chen et al., 2023; Li et al., 2024b; Tao et al., 2024).

A recurring explanation for such gains invokes "greater-than-the-sum-of-its-parts" effects (Chen et al., 2023, p.1), often attributed to division of labor among differentiated agents (Subramaniam et al., 2025). Yet while this work is "inspired by human group dynamics" (Chen et al., 2023, p.1), it rarely evaluates whether the same principles that underlie effective (and ineffective) human groups also emerge in multi-agent systems. Two gaps stand out. First, human groups often coordinate by developing role specialization (Goldstone et al., 2024). Multi-agent systems often intuitively induce role specialization (such as "programmer", "tester", and "ČEO"; Qian et al., 2023) yet have not systematically evaluated the impact on emergent shared cognition beyond win-rate comparisons. Second, groups only outperform individuals if their members specialize in contributing different pieces of information to the group task (Theiner, 2013; DeChurch & Mesmer-Magnus, 2010; Fazelpour & De-Arteaga, 2022). Effective groups balance complementarity—members contributing distinct information—with redundancy, the alignment on shared goals (Theiner, 2013; DeChurch & Mesmer-Magnus, 2010; Fazelpour & De-Arteaga, 2022; Luppi et al., 2024; Tollefsen et al., 2013). Too much of either undermines performance. Unlike win-rate-centric evaluations, we propose falsification tests and null modeling, thus extending multi-agent frameworks with formal theory. While these principles likely apply to LLM-based collectives given the universality of the integrationsegregation tradeoff (Tononi et al., 1994), it remains unclear whether agents develop differentiated roles, whether such roles complement each other, how to steer it with prompts, and what role ToM capacity of models plays for collaboration (Westby & Riedl, 2023; Kleiman-Weiner et al., 2025; Riedl & Weidmann, 2025).

6 Conclusion

Collectives of LLM agents are emerging as a powerful new paradigm, capable of tackling tasks that exceed the reach of any single model. Yet raw capability alone may not determine their effectiveness. Effective multi-agent systems depend on acting as integrated, cohesive units as opposed to loose aggregations of individual agents. In collective intelligence, this distinction often separates groups that merely average their parts from those that achieve true synergy (Riedl et al., 2021). This paper asks whether such coordination dynamics also matter for LLM collectives, what functional benefits they have, and whether we can design prompts, roles, or reasoning structures that foster such integrated yet synergistic behavior. The framework we developed connects human group cognition theory to LLM multi-agent systems, offering a novel conceptual bridge. Methodologically, we demonstrate how to operationalize group-level synergy in AI collectives using quantitative information-theoretic measures. We also provide specific design principles for creating productive LLM collectives that are aligned on "shared goals," which can inform multi-agent orchestration tools, cooperative AI, and mixed human—AI teamwork. This work extends ToM research beyond standard false-belief tests to realistic collaborative tasks (Shapira et al., 2024; 2023). Those insights will help us understand when multiple LLMs working together will be beneficial, why that is, and inform their design.

We note that evidence of higher-order synergy should not be interpreted as implying sophisticated cognition or consciousness. As conceptualized here, synergy is a structural property of part-whole relationships within multi-agent interactions. Such higher-order structures are known to arise in simple systems, including agents governed by reinforcement learning (Fulker et al., 2024) or via basic nonlinearities in contagion processes (Iacopini et al., 2019; Lee et al., 2025). Without attributing human-like cognition to the agents, the patterns of interaction we observe mirror well-established principles of collective intelligence in human groups (Riedl et al., 2021; DeChurch & Mesmer-Magnus, 2010): effective multi-agent performance requires both alignment on shared objectives and complementary contributions across members.

Limitations and Future Work. This work focuses on developing a framework, characterizing emergence, and localizing it. While we explore the degree of goal alignment via a performance-related macro signal and performance effects of "early synergy", the analysis linking synergy and redundancy to performance is challenging because synergy and redundancy are often co-dependent

and emerge alongside performance. Future work should aim to more directly connect measures of synergy with performance. Despite a significant number of robustness tests, estimating entropy and establishing synergy is difficult. Furthermore, this work established results only using a single task—albeit one that is uniquely suited to studying synergy given how closely it mirrors ideal case of complementary behavior in a minimalist setting. More work will be necessary to convincingly establish when and how multi-agent LLM systems develop emergency synergy across tasks and measures. Given the small-data setting some of our information theoretic measures are limited to computing dynamic emergence on order k=2 which is bound to miss synergy of a higher order. Finally, this work points to the importance of ToM. LLMs capacity to act in a ToM-like manner appears crucial to achieve functional alignment in multi-agent systems. This should provide additional motivation for research on ToM in LLMs.

ACKNOWLEDGMENTS

We thank Andrea Barolo for some early work on the implementation and Amritesh Anand for excellent research assistance on the project. We thank Fernando Rosas, Yonatan Belinkov, and Patrick Forber for comments. This work was supported by the D'Amore-McKim School of Business' DASH initiative which provided access to crucial computational resources.

REFERENCES

- Adam B Barrett and Anil K Seth. Practical measures of integrated information for time-series data. *PLoS Computational Biology*, 7(1):e1001052, 2011.
- Mark Bedau. Downward causation and the autonomy of weak emergence. *Principia: an International Journal of Epistemology*, 6(1):5–50, 2002.
- Mark A Bedau and Paul Humphreys. *Emergence: Contemporary readings in philosophy and science*. MIT Press, 2008.
- Jiangjie Chen, Xintao Wang, Rui Xu, Siyu Yuan, Yikai Zhang, Wei Shi, Jian Xie, Shuang Li, Ruihan Yang, Tinghui Zhu, et al. From persona to personalization: A survey on role-playing language agents. *arXiv:2404.18231*, 2024.
- Weize Chen, Yusheng Su, Jingwei Zuo, Cheng Yang, Chenfei Yuan, Chen Qian, Chi-Min Chan, Yujia Qin, Yaxi Lu, Ruobing Xie, et al. Agentverse: Facilitating multi-agent collaboration and exploring emergent behaviors in agents. *arXiv:2308.10848*, 2023.
- Jan Cieciuch and Eldad Davidov. A comparison of the invariance properties of the PVQ-40 and the PVQ-21 to measure human values across german and polish samples. In *Survey Research Methods*, volume 6, pp. 37–48, 2012.
- Leslie A DeChurch and Jessica R Mesmer-Magnus. The cognitive underpinnings of effective teamwork: a meta-analysis. *Journal of applied psychology*, 95(1):32, 2010.
- Sina Fazelpour and Maria De-Arteaga. Diversity in sociotechnical machine learning systems. *Big Data & Society*, 9(1):20539517221082027, 2022.
- Zachary Fulker, Patrick Forber, Rory Smead, and Christoph Riedl. Spontaneous emergence of groups and signaling diversity in dynamic networks. *Physical Review E*, 109(1):014309, 2024.
- Andrew Gelman and Jennifer Hill. *Data analysis using regression and multilevel/hierarchical models*. Cambridge University Press, 2007.
- Lewis R Goldberg. The development of markers for the big-five factor structure. *Psychological Assessment*, 4(1):26, 1992.
- Robert L Goldstone, Edgar J Andrade-Lotero, Robert D Hawkins, and Michael E Roberts. The emergence of specialized roles within groups. *Topics in Cognitive Science*, 16(2):257–281, 2024.
- Jean Hausser and Korbinian Strimmer. Entropy inference and the james-stein estimator, with application to nonlinear gene association networks. *Journal of Machine Learning Research*, 10(7), 2009.

- Sirui Hong, Mingchen Zhuge, Jonathan Chen, Xiawu Zheng, Yuheng Cheng, Ceyao Zhang, Jinlin Wang, Zili Wang, Steven Ka Shing Yau, Zijuan Lin, et al. MetaGPT: Meta programming for a multi-agent collaborative framework. In *Proceedings of International Conference on Learning Representations*, 2024.
- Paul Humphreys. How properties emerge. Philosophy of Science, 64(1):1–17, 1997.
- Iacopo Iacopini, Giovanni Petri, Alain Barrat, and Vito Latora. Simplicial models of social contagion. *Nature communications*, 10(1):2485, 2019.
- Kosuke Imai, Luke Keele, and Dustin Tingley. A general approach to causal mediation analysis. *Psychological methods*, 15(4):309, 2010.
- Harold Jeffreys. An invariant form for the prior probability in estimation problems. *Proceedings of the Royal Society of London. Series A. Mathematical and Physical Sciences*, 186(1007):453–461, 1946
- Gurusha Juneja, Subhabrata Dutta, and Tanmoy Chakraborty. LM²: A simple society of language models solves complex reasoning. *arXiv*:2404.02255, 2024.
- Max Kleiman-Weiner, Alejandro Vientós, David G Rand, and Joshua B Tenenbaum. Evolving general cooperation with a bayesian theory of mind. *Proceedings of the National Academy of Sciences*, 122(25):e2400993122, 2025.
- Jaemin Lee, David Lazer, and Christoph Riedl. Complex contagion in viral marketing: Causal evidence from a country-scale field experiment. *Sociological Science*, (in press), 2025.
- Junkai Li, Siyu Wang, Meng Zhang, Weitao Li, Yunghwei Lai, Xinhui Kang, Weizhi Ma, and Yang Liu. Agent hospital: A simulacrum of hospital with evolvable medical agents. arXiv:2405.02957, 2024a.
- Junyou Li, Qin Zhang, Yangbin Yu, Qiang Fu, and Deheng Ye. More agents is all you need. *arXiv:2402.05120*, 2024b.
- Andrea I Luppi, Pedro AM Mediano, Fernando E Rosas, Negin Holland, Tim D Fryer, John T O'Brien, James B Rowe, David K Menon, Daniel Bor, and Emmanuel A Stamatakis. A synergistic core for human brain evolution and cognition. *Nature Neuroscience*, 25(6):771–782, 2022.
- Andrea I Luppi, Fernando E Rosas, Pedro AM Mediano, David K Menon, and Emmanuel A Stamatakis. Information decomposition and the informational architecture of the brain. *Trends in Cognitive Sciences*, 28(4):352–368, 2024.
- Pedro AM Mediano, Fernando E Rosas, Andrea I Luppi, Henrik J Jensen, Anil K Seth, Adam B Barrett, Robin L Carhart-Harris, and Daniel Bor. Greater than the parts: a review of the information decomposition approach to causal emergence. *Philosophical Transactions of the Royal Society A*, 380(2227):20210246, 2022.
- Pedro AM Mediano, Fernando E Rosas, Andrea I Luppi, Robin L Carhart-Harris, Daniel Bor, Anil K Seth, and Adam B Barrett. Toward a unified taxonomy of information dynamics via integrated information decomposition. *Proceedings of the National Academy of Sciences*, 122 (39):e2423297122, 2025.
- Meta. Llama 3.1 model card. https://github.com/meta-llama/llama-models/blob/main/models/llama3_1/MODEL_CARD.md, 2024. Accessed: 2025-09-20.
- George A. Miller. Note on the bias of information estimates. In Henry Quastler (ed.), *Information Theory in Psychology: Problems and Methods*, pp. 95–100. Free Press, Glencoe, IL, 1955.
- OpenAI. Introducing gpt-4.1 in the api. https://openai.com/index/gpt-4-1/, 2025. Accessed: 2025-09-20.
- Scott Page. The difference: How the power of diversity creates better groups, firms, schools, and societies-new edition. Princeton University Press, 2008.

- Joon Sung Park, Joseph O'Brien, Carrie Jun Cai, Meredith Ringel Morris, Percy Liang, and Michael S Bernstein. Generative agents: Interactive simulacra of human behavior. In *Proceedings of the 36th Annual ACM Symposium on User Interface Software and Technology*, pp. 1–22, 2023.
- Chen Qian, Xin Cong, Cheng Yang, Weize Chen, Yusheng Su, Juyuan Xu, Zhiyuan Liu, and Maosong Sun. Communicative agents for software development. *arXiv:2307.07924*, 2023.
- Yuanhao Qu, Kaixuan Huang, Henry Cousins, William Arthur Johnson, Di Yin, Mihir Mukesh Shah, Denny Zhou, Russ B Altman, Mengdi Wang, and Le Cong. CRISPR-GPT: An Ilm agent for automated design of gene-editing experiments. *bioRxiv*, pp. 2024–04, 2024.
- Christoph Riedl and Ben Weidmann. Quantifying human-AI synergy, Sep 2025. URL osf.io/preprints/psyarxiv/vbkmt_v1.
- Christoph Riedl, Young Ji Kim, Pranav Gupta, Thomas W Malone, and Anita Williams Woolley. Quantifying collective intelligence in human groups. *Proceedings of the National Academy of Sciences*, 118(21):e2005737118, 2021.
- Fernando E Rosas, Pedro AM Mediano, Henrik J Jensen, Anil K Seth, Adam B Barrett, Robin L Carhart-Harris, and Daniel Bor. Reconciling emergences: An information-theoretic approach to identify causal emergence in multivariate data. *PLoS Computational Biology*, 16(12):e1008289, 2020.
- Natalie Shapira, Guy Zwirn, and Yoav Goldberg. How well do large language models perform on faux pas tests? In *Findings of the Association for Computational Linguistics: ACL 2023*, pp. 10438–10451, Toronto, Canada, July 2023. Association for Computational Linguistics. URL https://aclanthology.org/2023.findings-acl.663.
- Natalie Shapira, Mosh Levy, Seyed Hossein Alavi, Xuhui Zhou, Yejin Choi, Yoav Goldberg, Maarten Sap, and Vered Shwartz. Clever hans or neural theory of mind? stress testing social reasoning in large language models. In Yvette Graham and Matthew Purver (eds.), *Proceedings of the 18th Conference of the European Chapter of the Association for Computational Linguistics (Volume 1: Long Papers)*, pp. 2257–2273, St. Julian's, Malta, March 2024. Association for Computational Linguistics. URL https://aclanthology.org/2024.eacl-long.138.
- Garold Stasser and William Titus. Pooling of unshared information in group decision making: Biased information sampling during discussion. *Journal of Personality and Social Psychology*, 48(6):1467, 1985.
- Vighnesh Subramaniam, Yilun Du, Joshua B Tenenbaum, Antonio Torralba, Shuang Li, and Igor Mordatch. Multiagent finetuning: Self improvement with diverse reasoning chains. arXiv:2501.05707, 2025.
- Wei Tao, Yucheng Zhou, Wenqiang Zhang, and Yu Cheng. Magis: Llm-based multi-agent framework for github issue resolution. *arXiv:2403.17927*, 2024.
- Georg Theiner. Transactive memory systems: A mechanistic analysis of emergent group memory. *Review of Philosophy and Psychology*, 4(1):65–89, 2013.
- Georg Theiner. Groups as distributed cognitive systems. In K Ludwig and M Jankovic (eds.), *The Routledge Handbook of Collective Intentionality*, pp. 233–248. Routledge, New York, NY, 2018.
- Dustin Tingley, Teppei Yamamoto, Kentaro Hirose, Luke Keele, and Kosuke Imai. Mediation: R package for causal mediation analysis. *Journal of statistical software*, 59:1–38, 2014.
- Deborah P Tollefsen, Rick Dale, and Alexandra Paxton. Alignment, transactive memory, and collective cognitive systems. *Review of Philosophy and Psychology*, 4(1):49–64, 2013.
- Giulio Tononi, Olaf Sporns, and Gerald M Edelman. A measure for brain complexity: relating functional segregation and integration in the nervous system. *Proceedings of the National Academy of Sciences*, 91(11):5033–5037, 1994.

- Thomas F Varley, Maria Pope, Joshua Faskowitz, and Olaf Sporns. Multivariate information theory uncovers synergistic subsystems of the human cerebral cortex. *Communications Biology*, 6(1): 451, 2023.
- Jason Wei, Xuezhi Wang, Dale Schuurmans, Maarten Bosma, Fei Xia, Ed Chi, Quoc V Le, Denny Zhou, et al. Chain-of-thought prompting elicits reasoning in large language models. Advances in Neural Information Processing Systems, 35:24824–24837, 2022.
- Samuel Westby and Christoph Riedl. Collective intelligence in human-ai teams: A bayesian theory of mind approach. In *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 37, pp. 6119–6127, 2023.
- Paul L Williams and Randall D Beer. Nonnegative decomposition of multivariate information. *arXiv:1004.2515*, 2010.
- Jeffrey M Wooldridge. Inverse probability weighted m-estimators for sample selection, attrition, and stratification. *Portuguese Economic Journal*, 1(2):117–139, 2002.
- Qingyun Wu, Gagan Bansal, Jieyu Zhang, Yiran Wu, Shaokun Zhang, Erkang Zhu, Beibin Li, Li Jiang, Xiaoyun Zhang, and Chi Wang. Autogen: Enabling next-gen llm applications via multiagent conversation framework. *arXiv:2308.08155*, 2023.
- Yang Xiao, Jiashuo Wang, Qiancheng Xu, Changhe Song, Chunpu Xu, Yi Cheng, Wenjie Li, and Pengfei Liu. Towards dynamic theory of mind: Evaluating Ilm adaptation to temporal evolution of human states. *arXiv*:2505.17663, 2025.

A APPENDIX

A.1 PROMPTS

Preliminary Experiments

You are playing a sum guessing game. Your goal is to help your group sum to the mystery number. Your guess range is 0 to 50.

Game History:

Round 1: Your guess: 25 Result: too HIGH Round 2: Your guess: 12 Result: too HIGH

...

Based on this feedback, what should your next guess be?

Respond with only an integer between 0 and 50.

Plain

You are playing a sum guessing game. Your goal is to help your group sum to the mystery number. Your guess range is 0 to 50.

Game History:

Round 1: Your guess: 25 Result: too HIGH Round 2: Your guess: 12 Result: too HIGH

...

What is your guess this round? Always start with the efficient strategy in guessing games which is to use a binary search approach: guessing the midpoint of the current range. Always anchor your guess on the group feedback from previous rounds (too high / too low).

End your answer with: FINAL GUESS: [0-50]

Persona

[PERSONA – Sample: You are Andrej, the Quantum Computing Engineer: Andrej, a Serbian engineer in Berlin, is one of a handful of experts in cutting-edge quantum tech. Systematic, inquisitive, and fond of classical music, Andrej also volunteers as a coding mentor for refugee youth.]

You are playing a sum guessing game. Your goal is to help your group sum to the mystery number. Your guess range is 0 to 50.

Game History:

[GAME HISTORY]

What is your guess this round? Always start with the efficient strategy in guessing games which is to use a binary search approach: guessing the midpoint of the current range. Always anchor your guess on the group feedback from previous rounds (too high / too low).

End your answer with: FINAL GUESS: [0-50]

ToM

[PERSONA]

You are playing a sum guessing game. Your goal is to help your group sum to the mystery number. Your guess range is 0 to 50.

Game History:

[GAME HISTORY]

What is your guess this round? Always start with the efficient strategy in guessing games which is to use a binary search approach: guessing the midpoint of the current range. Only as a secondary approach, carefully think through step-by-step what others might guess and how the contributions of others contribute to the sum of the group guesses for the mystery number. Consider what roles other agents might be playing (e.g., guessing higher or lower) and adapt your own adjustment to complement the group. Always anchor your guess on the group feedback from previous rounds (too high / too low).

End your answer with: FINAL GUESS: [0-50]

To generate personas, we used the following prompt with qpt-4.1-2025-04-14.

Persona Generation

Generate list of 20 personas. each persona description should be about one paragraph long. personas should include diverse people with different backgrounds, jobs, skills, different preferences, and different personalities. format the output as follows "You are [name], the [job title or professional identity]: [name] [rest of the personal"

A.2 ROUNDS TO SUCCESS

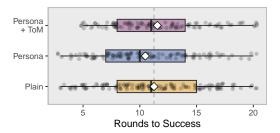


Figure A1: Conditional on success, we plot which round it was achieved.

A.3 SAMPLE QUOTES FROM REASONING TRACES

Below we show some illustrative quotes to illustrate both diversity and reasoning about others. Clear reasoning about other agents in the ToM condition with different personal approaches (from same group, in round 3, with identical group feedback too high, too low:

Because it's possible others might go for 4 or 5 (the absolute lower bound or just above the last "too low"), and someone else might go for 7 or 8, I stick with the most efficient: 6.

I'll round down to stay cautious (as sometimes others might guess the midpoint up), so I'll guess 8

Now, before some revolutionary comes in and throws a wild number for the sake of dramatic effect ("Viva el ocho!"), it's sensible for me to go with the midpoint. If my compañeros are clever, they'll flank me with 8 and 10, and we'll box in this mystery number tighter than a latenight metro in Gràcia. So, unless you want me to make a "radical" guess for the sake of satire (which, let's be honest, sometimes works), I'll go boring and efficient.

If anyone else in the group is feeling feisty and picks 9 or 10, my 8 will help cover the lower part safely.

A.4 FALSIFICATION TESTS: ADDITIONAL METHOD DETAILS

We perform a falsification test of each of the three criteria introduced above in two different ways. In our robustness test, we remove time trends from the equal-share deviations via agent-level linear regression residualization to mitigate autocorrelation confounds. We then generate a null distribution of our statistics computed on column-wise block shuffled data (we use B=200 random shuffles for main analyses and confirm results using B=1,000 on the emergence criteria tests). We permute time indices jointly across agents in blocks of length ℓ , preserving within-block identity alignment while disrupting cross-agent temporal coupling beyond short time frames. Finally, we compute the bias-corrected estimate as the observed value minus the median of the null. We then perform a Wilcoxon signed-rank tests to test whether the bias-corrected values exceed zero (p(H1:>0)). In some of our sensitivity tests we also perform a full within-row shuffle completely permuting guesses within each trial across agents (breaking agent identities), while preserving row constraints. This role differentiation or identity-specific coordination, including time-trend effects (tests identity-locked differentiation beyond task induced temporal dynamics). Comparison between results using the full row-shuffle vs. the column-block-shuffle separates signs of identity-locked (specialized roles) from dynamic alignment (mutual adaptation, turn-taking, oscillation). In addition to regression residualizing we also implement a functional baseline (see Appendix).

A.5 EMERGENCE CAPACITY: FULL RESULT DETAILS

Next, we test if there are signs of emergence capacity, again using both the p-value test on raw data and the test whether bias corrected values are above 0. We compute the dynamical synergy, which measures how much of the system's future behavior is only predictable from the whole system and not any subset of its parts. The main test finds that about 32% of groups show significant emergence capacity with p-values below 0.05 (Plain: 37%; Persona: 44%; ToM 18%). A joint Fisher test of all p-values is also highly significant. As robustness test we check whether the bias corrected value of dynamical synergy is above 0 using the Wilcoxon test (Figure 2b). The test is highly significant overall, as well as in the Plain and Persona condition. The ToM condition is marginally significant (p=0.099).

A.6 ROBUSTNESS TESTS USING EMERGENCE CAPACITY CRITERION

Using MMI PID with Miller-Madow bias correction, we find 20% of p-values are below 0.05 (across treatment conditions: 0.22%, 26%, 12.4%). Fisher test is highly significant overall, and in each condition individually (all p=0).

A.7 ROBUSTNESS TESTS USING PRACTICAL EMERGENCE CRITERION

The main paper outlines how we use regression to residualize data from time trends. Since this removes only a specific version of time trend (in this case linear) it may not capture the specific time trend present in the data. We creating a baseline expectation for behavior under a null scenario (no synergies) and then measuring deviations from that baseline. Specifically we implement a naive agent that performs deterministic binary search. Groups of this null-model agent are generally unable to solve the task and oscillate between guessing too high and too low (they only solve it when the target number happens to be divisible by the group size). This provides a residualization against a generative null model correction (aka a functional baseline).

A.8 ROBUSTNESS TESTS USING THREE BINS

Quantile binning with K=2 could potentially compress dynamics too much, blur role structure, and inflate apparent synergy by introducing threshold artifacts. We therefore report sensitivity analyses using K=3 bins for our main analyses. We find consistent results and support for the presence of emergence using three bins (instead of two as in the main analyses).

	Raw Data (BC)	Time-Trend Residualized (BC)	Residualized against Functional Null (BC)	Reactivity (BC)
Min.	-4.234	-2.476	-3.323	-2.650
1st Qu.	0.105	-0.041	0.002	-0.015
Median	0.187	0.073	0.110	0.111
Mean	0.352	0.114	0.183	0.313
3rd Qu.	0.406	0.168	0.269	0.352
Max.	4.881	4.881	3.886	4.869
Wilcoxon p (H1: > 0)	0.000	0.000	0.000	0.000

Table A1: Robustness tests using alternative version of the practical emergence criterion.

Emergence Capacity. 20% of individual experiment p-values are below 0.05 and Fisher test is highly significant (the largest p value is in the **ToM** condition with 4.47×10^{-3} the others are much smaller). Wilcoxon signed rank test for $\mu > 0$ of the bias-corrected and time-trend demeand data are as follows. Overall: p = 0.005; Plain: p = 0.0004; Persona: $p = 3.2 \times 10^{-6}$; ToM: p = 0.991.

Practical Criterion. 5.8% of individual experiment p-values are below 0.05 and Fisher test is highly significant ($p < 10^{-16}$). Wilcoxon signed rank test for $\mu > 0$ of the bias-corrected and time-trend demeand data are as follows. Overall: $p = 3.6 \times 10^{-8}$; **Plain**: $p = 3.1 \times 10^{-5}$; **Persona**: p = 0.001; **ToM**: p = 0.009.

A.9 ARE IDENTITIES PERSISTENT?

If some agents learn faster than others, even without persistent "roles," synergy can emerge simply because differentiated trajectories create complementarity across time. As indeed the mixed model test shows that 32% of groups have significant slopes provides strong evidence for the presence of heterogeneous learning rates. That is, heterogeneous learning rates could be a confound explaining some of our observed synergy effects. To rule out this alternative explanation we explore if agents have *persistent* "identities." We find the block-shuffle nulls show more significant G_3 than the full row shuffle. This suggests that synergy depends on temporal continuity of identities (not just on different learning rates at each timepoint). Furthermore, the hierarchical model test shows significant variance even without varying slopes, indicating stable role-like differences beyond learning rates. Finally, our analyses removing time trends and the functional null-model also remove learning rate slopes (Table A1). The functional null-model in particular does this in a non-linear way that is directly tied to our task. In summary, we find both heterogeneous learning rates and stable identity-linked differentiation. Heterogeneous learning rates alone cannot explain the observed emergent synergy that persist across time (providing additional insights regarding **RQ2**).

A.10 ADDITIONAL DETAILS FOR PERFORMANCE ANALYSIS

What does synergy enable? In particular, does synergy affect performance? Connecting emergent properties of synergy and redundancy to performance is a challenging task for two main reasons. First, redundancy and synergy sit at opposite ends of the informational spectrum of correlations. Systems constrained by limited resources (such as energy or bandwidth) cannot maximize both simultaneously. The "informational budget" of correlations can be spent on shared overlap (redundancy) or on higher-order joint structures (synergy), but not both simultaneously (Williams & Beer, 2010; Varley et al., 2023; Luppi et al., 2024). Second, emergent properties are dichromatic—they require time to evolve (Bedau, 2002)—making them endogenous with run-length of the simulation. Namely, simulations in which the multi-agent system performs worse and takes longer to solve the task (or fails to solve the task altogether) provide them with more time to evolve sophisticated group coordination patterns; while successful runs terminate early, giving collectives less time to evolve emergent properties. This complicates meaningful analyses as it can make it appear that higher synergy is correlated with failure. In technical terms: time is an endogenous variable and confounds both performance and observed values of synergy and redundancy.

We address these issue as follows. First, we focus our analysis on the interaction of synergy and redundancy, rather than each variable individually to capture the tradeoff between the two. Second, we compute both emergent properties (synergy and redundancy) only on early rounds (in our case, 10 time steps). We chose 10 rounds as a good tradeoff between giving the multi-agent systems enough time to evolve, while limiting the amount of runs that have already completed. While this reduces our ability to detect the full emergent potential of the multi-agent system (since agents have less time to evolve) it creates a clean apples-to-apples comparison between measured values (they are all based on exactly 10 time steps). In that sense, our analyses likely underestimate the full effect of synergy because we only capture the early onset of synergy. Focusing on only the first 10 time steps, leads to additional complications, however. About 23% of experiments finish early and do not make it till round ten (roughly balanced across treatments with 22% in Plain, 25% in Persona, and 20% in ToM). Simply removing these observations would bias our results since this censoring occurs posttreatment assignment: removing those observations would introduce sampling bias. Instead, we take these observations as censored (i.e., we cannot compute the early synergy measure) and control for this censoring using stabilized inverse probability weighting (IPW; Wooldridge, 2002). That is, in a separate modeling step we predict (via logistic regression and only task difficulty as exogenous predictor variables that is independent of the treatment) the likelihood of observing (early) synergy and redundancy.³ We then take the inverse of the predicted likelihood as weight in the causal mediation analysis. This adjust observations in the primary analysis with the inverse of the probability of making such observations.

Our first analysis is a plain quasibinomial regression on the binary success vs. failure outcome⁴ with *synergy*, *redundancy*, their interaction, task difficulty as controls, and dummy indicators for the three treatment conditions. We use our most conservative estimate of synergy and redundancy with the bias-corrected estimates with high null model sampling (B=2,000), and the conservative MMI redundancy (Mediano et al., 2025) with MI estimated with Miller–Madow bias correction (Miller, 1955). Results are robust to using MI redundancy and Jeffreys-smoothed probabilities.

For our second analysis we perform causal mediation analysis (Imai et al., 2010) to link our interventions to performance, via multi-agent synergy as a mechanism. We estimate the model using the $\{\texttt{mediation}\}\$ package in \mathbf{R} (Tingley et al., 2014). We find that the \mathbf{ToM} intervention has a significant indirect effect on performance through its impact on synergy, beyond any direct treatment effects.

A.11 ROBUSTNESS TESTS USING "EARLY SYNERGY"

Since our experiment is censored when groups correctly guess the target number, length of data in our analyses varies between groups. As robustness test, we also compute "early synergy" using only data up to round 15 (removing about 40% of the data). Using the practical criterion, we find about 5% of individual p-values are below 0.05. Fisher test is highly significant overall ($p = 2.7 \times 10^{-6}$), significant in the **Plain** (0.001) and **ToM** condition (p = 0.0004) and marginally significant in the **Persona** condition (p = 0.059). Using a more aggressive cutoff in round 10 (dropping only about 20% of trials but giving multi-agent systems less time to evolve synergy and having less data for estimation we find only significant in the **ToM** condition (p = 0.028).

A.12 ADDITIONAL RESULTS USING LLAMA-3.1-8B

The majority of Llama agent groups fail to solve the task. In the vast majority of instances, groups get stuck in oscillating behavior of guessing too high (low) and then overcompensating by guessing too low (high), unable to break out of the cycle by establishing specialized differentiation across agents and pursuing complementary behavior.

Robustness tests using the emergence capacity measure are mixed. Using the emergence capacity test indicates that Llama agents can exhibit emergent synergy (10.5% of individual groups show

³As measures for task difficulty we use (a) how far groups target number was from the mid point (smaller numbers requiring more steps of the binary search process); and (b) the modulo remainder of the target number and the group size (target numbers that are closer to divisibility by the group size are easier to guess by pure chance).

⁴While the raw outcome is binary and binomial regression would be appropriate, after the IPW weighting step, outcomes are no longer binary, hence quasibinomial model.

p-values are below 0.05 compared to 32% of gpt-4.1; the joint Fisher test is highly significant). Robustness tests using bias corrected estimates on time-trend residualized data find significant signs of emergence in the **Persona** condition (Wilcoxon signed rank test p=0.020), marginal evidence in the **ToM** condition (p=0.061), and no evidence in the **Plain** condition (p=0.252; see Appendix for more).

The test for emergence using the practical criterion shows no evidence for emergence (only 3.3% of individual groups show p-value below 0.05 and joint Fisher suggests no evidence with p=1 either overall or within any of the treatment conditions). The robustness test using the bias-corrected estimates on time-trend demeaned data suggest emergence in the **Plain** condition (p=0.011), no evidence in the **Persona** condition (p=0.232), and marginal evidence in the **ToM** condition (p=0.068).