
Prompting as Scientific Inquiry

Ari Holtzman

Department of Computer Science
University of Chicago
Chicago, IL, 60637
aholtzman@uchicago.edu

Chenhao Tan

Department of Computer Science
University of Chicago
Chicago, IL, 60637
chenhao@uchicago.edu

Abstract

Prompting is the primary method by which we study and control large language models. It is also one of the most powerful: nearly every major capability attributed to LLMs—few-shot learning, chain-of-thought, constitutional AI—was first unlocked through prompting. Yet prompting is rarely treated as science and is frequently frowned upon as alchemy. We argue that this is a category error. If we treat LLMs as a new kind of complex and opaque organism that is trained rather than programmed, then prompting is not a workaround: it is behavioral science. Mechanistic interpretability peers into the neural substrate, prompting probes the model in its native interface: language. We contend that prompting is not inferior, but rather a key component in the science of LLMs.

1 Introduction

Olah et al. (2020) make the case for mechanistic interpretability as the study of a new kind of system:

In this view, neural networks are an object of empirical investigation, perhaps similar to an organism in biology. Such work would try to make empirical claims about a given network, which could be held to the standard of falsifiability.

Of course, nothing in this description is specific to mechanistic interpretability. Indeed, many of the most profound discoveries about how LLMs work—in-context learning (Brown et al., 2020), chain-of-thought prompting (Wei et al., 2022), the reversal curse (Berglund et al., 2023)—have used a model’s natural inputs and outputs for stimuli and response.

If we discovered an intelligent alien species, we would learn many things from playing simple card games with them and observing their reasoning patterns that we would have trouble learning by dissection. Similarly, prompting allows us to probe LLMs through their optimized communication channel, revealing capabilities and limitations that might remain hidden in their weights and activations if we don’t know what to look for.

However, prompting is frequently considered inferior to other methods, especially those used in Mechanistic Interpretability that attempt to explain behavior at the level of neurons or individual parameters. When discussing methods for improving models, researchers commonly dismiss methods of structured prompting, treating it as a mere hack rather than genuine progress. Until, of course, such a method is endowed with a name and enough proven success so that it is viewed as distant from the word “prompting”, e.g., In-Context Learning (Brown et al., 2020).

Why the bias against prompting? In this position paper we argue that *prompt science*—performing experiments that use a model’s natural inputs and outputs for intervention and measurement—has been conflated with *prompt engineering*—the arduous process of finding the best prompt for a specific model, task, and dataset, often by brute-force search or by using techniques that are poorly understood without attempting to explain why they work (Figure 1 provides an overview of our key arguments).

Preprint.

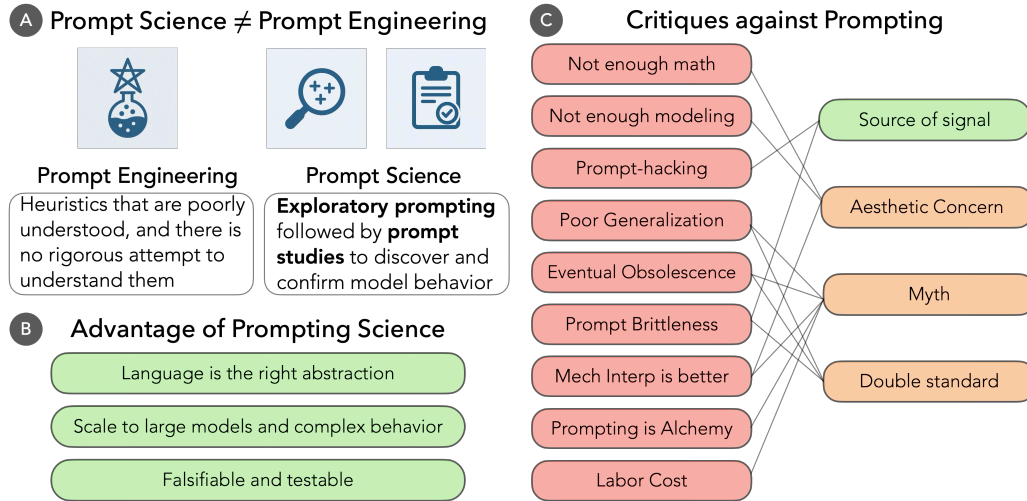


Figure 1: An overview of key arguments in this work. A. It is important to distinguish *prompt science* from *prompt engineering*. B. Prompt science enables studies at the right abstraction, scales to large models and complex behavior, and is falsifiable and testable. C. We point out weaknesses of typical critiques against prompting and provide detailed rebuttal in Section 6.

Similar complaints have been made about both feature-engineering and hyperparameter-tuning—that they are often arbitrary, opaque, and take labor and compute that could be better spent elsewhere. As an empirical field, some amount of “fiddling for better performance” is necessary to verify the robustness of our results, but there is no reason to equate this fiddling with studies that attempt to understand *how* these parameters affect the systems they are embedded in.

Using prompts to understand, discover, and control model behavior, *prompting as scientific inquiry*, remains the most impactful method we have for finding out what lies inside LLMs since the leap from language models to large language models and *prompting studies* are usually how we experimentally verify these behaviors. Training innovations have made these models better, more efficient, and capable of handling longer contexts, but it has generally been difficult to explore what improved capabilities these advancements offer through optimization methods alone. Instead, we have relied on the richness of language, with all the subproblems it can express, to prompt LLMs into revealing their capabilities and limitations.

The bias against prompting is part of our field’s persistent preference for optimization over exploration, for mathematical elegance over empirical discovery. Yet history repeatedly demonstrates that prompting is the source of our most significant insights. In-Context Learning wasn’t the projected outcome of scaling-up Transformers, it was a delightful side-effect that changed the game of what we could expect from model generalization at inference time. Unsupervised chain-of-thought reasoning wasn’t discovered through architecture innovation but by simply asking models to “think step by step.” Many alignment techniques are, at their core, elaborate prompting methods dressed in mathematical formalism to appeal to the aesthetics of rigor in ML communities.

Behind the hand-waving dismissals lies an uncomfortable truth: prompting reveals model capabilities we never knew to look for, while interpretability has, so far, largely confirmed hypotheses we’ve already formulated. Prompting is a form of exploratory science that is badly needed to explore the complex and opaque new organism we have been confronted with: Large Language Models. The scientific community’s fascination with “opening the box” has often overshadowed the insights available through more accessible means. Because we *can* inspect weights, neurons, and attention patterns numerically, we *do*—even in cases when those investigations continually yield little actionable understanding. If we only had prompting at our disposal—if LLMs were truly black boxes—we would likely have made greater strides in probing their behavior systematically, designing experiments with rigor and creativity to uncover their internal structure through inputs and outputs alone. Even without whitebox access, LLMs are easier to study than people: LLMs do not suffer from observer effects, can be made (almost) deterministic, and allow us to explore many experiments on the same subject in parallel.

In this paper, we advocate for exploratory prompting as an act of scientific inquiry and prompting studies as key means of confirming hypotheses. We start by highlighting the impact that prompting has led to in the recent advances of large language models (Section 2), and then provide our account of how prompting got counted out (Section 3). We argue that prompting should be a critical form of scientific inquiry for interrogating LLMs (Section 4) by 1) providing a direct comparison between prompting and mechanistic interpretability to illustrate that they are complementary and critical methods to understand and control large language models (Section 5), and 2) giving direct rebuttal to counterarguments in Section 6.¹ Finally, we lay out concrete future directions to pursue prompting as scientific inquiry (Section 7).

2 Prompt Engineering vs. Prompt Science

In this work, we define *prompting* as interacting with LLMs with natural language and observing their behavior, including the output texts and/or the probability distributions. We would like to distinguish between *prompt engineering* and *prompt science*.

- **Prompt Engineering** is the process of optimizing the prompt for a specific model, task, and dataset, often by brute-force search or heuristics techniques that are poorly understood and that the engineer does not attempt to understand.
- **Prompt Science** is the process of using the inputs of LLMs to discover and confirm regularities in their output distributions. *Exploratory prompting* means probing the model with new prompts to discover new behavior while *prompting studies* are formal behavioral studies that use vary prompts in structured ways to confirm a hypothesis.

We believe that much of the bias against prompt science, comes from its association with prompt engineering, which does not have scientific goals. Importantly *prompt science* must make falsifiable claims (e.g., adding “let’s think step by step” to the end of a prompt improves accuracy on math benchmarks), whereas prompt engineering generally just optimizes prompts for a certain situation without making claims about the prompts produced.

But prompting is not just a way to use models, it is the primary means by which we have discovered what they can do. Across model generations, discovery via prompting has consistently outpaced our understanding of the internal structure of LLMs, surfacing capabilities that would have otherwise gone unnoticed. Importantly, these discoveries were rarely the result of theoretical foresight or interpretability breakthroughs. They came from structured interaction: trying things out, observing behavior, and refining prompts until something new emerged. Many people would describe this as unscientific. We disagree: this is the essence of scientific discovery—intervening on key variables in a system is how we discover new effects to test in the first place. It just so happens that the inputs to this system is human language.

Below, we present three case studies of this dynamic. Each represents a major leap in how we understand and control LLMs—one that began not with weight analysis, but with well-chosen words.

Case Study: Sparks and Embers. An essential issue that researchers must grapple with regarding LLMs is the demand to measure what new capabilities models may have, by finding instances of novel behavior. Claims of novel behavior often get debunked when behavior does not generalize as one might hope—just because a model can solve 1 digit addition problems does not mean it can solve three digit ones. This is an issue that psychologists run into as well (Yarkoni, 2022), but it is one that the science of LLMs has better machinery to deal with: we can go back to the same subject and perturb the experiment to see if a certain difference would have mattered, a luxury that human studies do not have. Recently “Sparks of Artificial Generative Intelligence” (Bubeck et al., 2023) suggests new emergent capabilities some of which were shown to lack robustness by “Embers of Autoregression” (McCoy et al., 2024). **Both** of these are prompting studies, operating at the black-box level through carefully structured input text. Bubeck et al. (2023) is an exploratory study of what an early version of GPT-4, which suggested a number of new capabilities such as the ability to generate stories with embedded hidden messages (via acrostics) or write a small 3D game from scratch, discovered through exploratory prompting. Yet McCoy et al. (2024) showed

¹To make the argument non-trivial, we focus on the case of understanding, developing, and controlling LLMs. It is already very clear that prompting is the essential ingredient in the emerging area of AI agents.

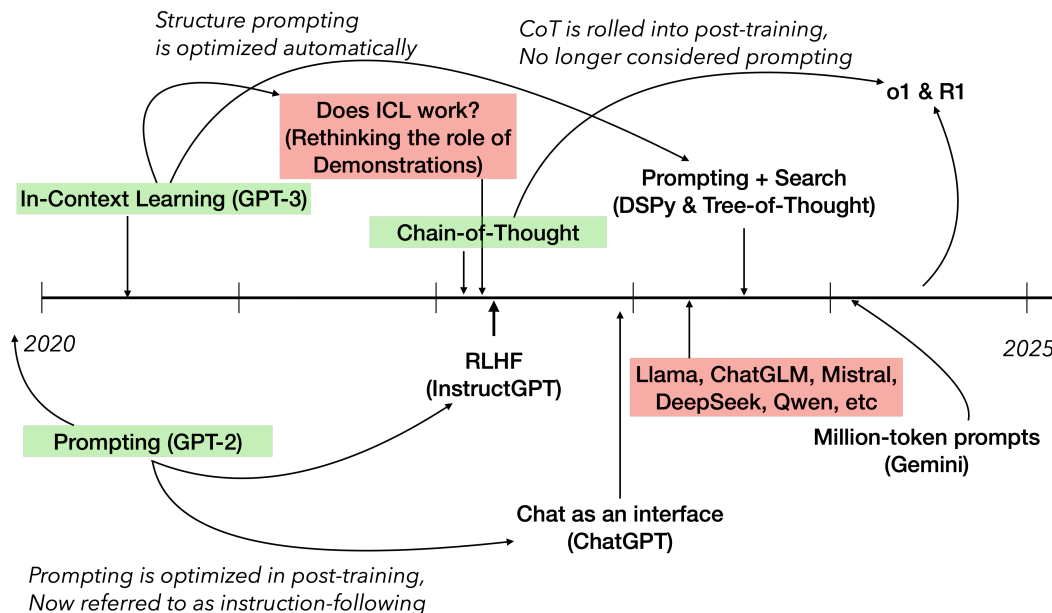


Figure 2: An illustration of the timeline related to the success of prompting (highlighted in green) and potential causes of the poor perception of prompting (highlighted in red).

some of these incredible abilities didn’t generalize in the way we’d hope, e.g., that the model can *decode* a shift cipher² but have trouble *encoding* that same cipher. This was shown through careful prompting studies: finding specific situations where an aspect of the prompt can be carefully varied (e.g., encoding vs. decoding), and looking at the change in decoded output. Prompting was used in these studies because it is the most natural interface under which to manipulate the conditional distribution of LLMs, both for discovery and hypothesis confirmation.

Case Study: Chain-of-Thought. Wei et al. (2022) showed that using in-context learning to generate human-like reasoning chains resulted in higher math scores for LLMs, while Kojima et al. (2022) built on the tradition of exploratory prompting popularized by GPT-2’s prompting studies (e.g., using “TL;DR” for generating summaries (Radford et al., 2019)), to show this could be done by simply appending five words to a prompt: “let’s think step-by-step.” The intuition was that strings used to signal when humans would generate reasoning chains could *prompt* models to enact this behavior. The key insight wasn’t that models simply needed more raw computational power or an method to engineer such computation into a model manually, though later studies showed that models could use inference-time compute without “thinking” out loud (Pfau et al., 2024). Researchers instead took advantage of behavioral patterns models learned from their training corpora, leveraging the correlations between reasoning chains and correct outputs. This discovery has since spawned a new generation of training procedures that leverage *inference-time compute* strategies. Modern systems like DeepSeek-R1 (Guo et al., 2025) and GPT-4o1 (OpenAI, 2024) now demonstrate that models can reason better in recognizable ways—achieving improved results on complex mathematical problems and challenging domain-specific tasks such as graduate-level questions in GPQA (Rein et al.). Asking models to show their work revealed how we could train models to do something more like thinking, a process of exploration and refinement that would have been difficult to achieve without prompting.

Case Study: Constitutional AI. Prompting is also the backbone of many alignment strategies. In Constitutional AI (Bai et al., 2022), prompting is used not only to elicit desirable behavior, but to enforce constraints and shape it. The process works by having models generate self-critiques and revisions, then fine-tuning the original model on these revised responses. Rather than relying entirely on human feedback, models are prompted to critique their own answers against a predefined set of normative principles (the “constitution”), and then to revise those answers accordingly. Human oversight is provided through a list of rules or principles expressed in natural language—it is this

²A shift cipher is one where the alphabet is treated like a ring and letters are “shifted” by a fixed amount, e.g., in Rot3 letters are shifted by 3, so $A \rightarrow D, B \rightarrow E, \dots, Z \rightarrow C$.

form of natural language prompting that governs the entire process. Here prompting serves as both a tool of supervision and instrument of control, a counterpoint to the belief many researchers have that models could not improve by training on their own outputs, instead insisting that these systems could only degrade via self-training. Like Chain-of-Thought, Constitutional AI was not born from theoretical introspection or circuit-level understanding, but from treating the model as a behavioral agent that could be instructed, guided, and corrected through language.

3 How did Prompting Get Counted Out?

Given the huge impact prompting has had on the evolution, use, and theory of LLMs, how did it become “dark magic”? We provide our interpretation of what has happened.

First, a series of studies demonstrated the sensitivity of prompting (Arora et al., 2022; Min et al., 2022; Wallace et al., 2019, *inter alia*). Model behavior can be unpredictable, counterintuitive, and small changes can change the performance of a system drastically. As a result, researcher assumed that prompting cannot be trustworthy or useful, and many suggested we need to find a way to get rid of it (Kosch & Feger, 2025; Meincke et al., 2025; Morris, 2024, *inter alia*).

Second, there is a strong bias in the machine learning community towards algorithms, modeling, and training methods (Birhane et al., 2022; Lipton & Steinhardt, 2019; Sculley et al., 2014, *inter alia*). This bias, combined with the increasing availability of open-weight models, caused researchers and practitioners to continually pursue bespoke systems with novel methods despite the fact that large models consistently outperformed them with proper prompting (Nori et al., 2023).

Consequently, a curious pattern has emerged: the moment prompting discoveries prove transformative, they are quietly rebranded and distanced from their origins, see Figure 2. Chain-of-Thought reasoning (Wei et al., 2022) was discovered via In-Context Learning and then simplified to only require five words: “let’s think step-by-step” (Kojima et al., 2022). Once this proved effective, it became “inference-time compute” and was rolled directly into training procedures for models like o1 and DeepSeek-R1 (Guo et al., 2025)—no longer considered “prompting.” RLHF (Ouyang et al., 2022) and ChatGPT (OpenAI, 2022) essentially automated effective prompting strategies into post-training, but were celebrated as training innovations rather than systematized prompting. Yet prompting is still the way we understand the limits of these new models. Unfortunately, the decline of prompting studies inadvertently limits our understanding of these transformed models.

The pattern continued across multiple fronts: structured prompt optimization in DSPy (Khattab et al., 2024) is touted as “programming—not prompting—Foundation Models”, despite being fundamentally about systematic prompt design. Tree-of-Thought (Yao et al., 2023) framed itself as “exploration over coherent units of text” rather than acknowledging its foundation in structured self-prompting strategies. Constitutional AI (Bai et al., 2022) and Self-Instruct (Wang et al., 2023), while genuinely innovative, positioned themselves as training methodologies that transcended prompting’s supposed unreliability—despite relying heavily on carefully crafted prompts for self-critique and instruction generation. The consistent pattern is that successful prompting methods get rebranded once they prove effective, distancing themselves from the “prompting” label while retaining its core mechanisms. This distracts from the shared methodology, and from the fact that we should be leaning-in to “language as a probe” when it comes to LLMs.

Meanwhile, the field developed longer contexts—Gemini’s million-token capacity (Team et al., 2024) enabling many-shot prompting (Anil et al., 2024)—precisely to accommodate more sophisticated prompting strategies. Yet these advances were celebrated as architectural achievements, even though *the driving motivation behind long-context models is to be able to prompt with more data*. However, we have more and more long-context training papers, but not a lot of exploration on prompting strategies with long context models. The uncomfortable truth remains: virtually every major capability breakthrough originated from careful prompting experiments, yet prompting itself continues to be dismissed as unscientific hackery.

4 Prompting as Scientific Inquiry

Prompting is useful for three reasons that mutually support each other: (1) humans have a deep understanding of the structure of language; (2) since language models were discovered not designed

(Holtzman et al., 2023), prompting is the most direct way to discover novel behavior in LLMs; (3) prompting, while not altering the model internals, is falsifiable and testable, especially with open-weight models. This makes prompting a crucial form of scientific inquiry for interrogating LLMs and the structure of their computations.

Other methods we have available (e.g., discovering locally linear decision boundaries, statistical testing, etc.) largely select hypotheses from a very constrained set or reject null hypotheses. What we want, is a way to discover new capabilities, without having to be overly prescriptive about the hypothesis set. More formally, we want to find behavior in LLMs that was *a priori* unlikely (because previous models couldn't do it), but *a posteriori* predictable (because it can be used in reliable ways, e.g., ICL or CoT).

Despite the fact that it is standard practice for LLM researchers prompt models to understand aspects of how they work, it has become mentally associated with prompt engineers attempting to find tricks that improve benchmark numbers but are not generalizable. Part of this is the low barrier to entry, since many who use prompting do not use it rigorously. Yet this is hardly an argument against using prompting rigorously.

Putting LLMs in new contexts, examining the generated outcomes, and attempting to find patterns in the resulting probability distributions, these are precisely the necessary ingredients of exploratory science. And exploratory science is vital when we are confronted with a new system that evades intuitive explanation, as we are today. A similar point was made eight years ago in Neuroscience, by a paper titled “Could a neuroscientist understand a microprocessor?” (Jonas & Kording, 2017). The answer is “No.” precisely because the strategies taken would not exploit the regularity with which these systems could be manipulated. LLMs have been handed to us on a silver platter, from this perspective: they already speak *approximately* the same language we do, shouldn't we be excited to be able to probe them with it?

Contrast this with a recently popular line of work in mechanistic interpretability: Sparse Autoencoders (SAEs) (Bricken et al., 2023; Cunningham et al., 2024; Gao et al., 2024, *inter alia*). What is the goal of SAEs? It is to discover concepts in models we did not know were there. This is very similar problem to the one prompting solves, but there is a catch: a plethora of evidence shows that they are unreliable (Makelov et al., 2024; Paulo & Belrose, 2025; Smith et al., 2025). Prompting allows us to probe for *what might be there* using a very natural UI, that of language. What prompting finds is also inherently testable: simple check if similar prompts work in different contexts or on different models.

Why is there such a perceived contrast between prompting and mechanistic interpretability? Because of the *mythos of mechanistic interpretability*, which suggests that if we work with the internal mechanisms of models our results are not subjective like they are with language. Seven years ago Lipton (2018) explained that “claims regarding interpretability of various models may exhibit a quasi-scientific character.” The underlying concept of interpretability remains ill-defined seven years later, warning us that the absence of clear definitions can create a false sense of scientific rigor where none actually exists. In the age of *mechanistic* interpretability we face a new but subtler issue: mechanistic interpretability has found ways of being more falsifiable than many previous interpretability methods (e.g., through interventional studies), but also **mechanistic interpretability has codified the idea that a rigorous statement about an LLM must be about model weights**.

This is simply aesthetic bias against blackbox analysis. The probability distributions emitted by LLMs are as much an objective part of their mechanistic description as anything else. Prompt science is falsifiable and scientifically rigorous. The main source of error is human presumption: humans may think a comma in a prompt is meaningless, but the model may have noticed that in the training set a correctly placed comma predicts an 80% change of answering a math problem correctly, but an incorrectly placed comma implies only 40% accuracy. This bias exists in mechanistic interpretability too, but is often subtler, e.g., the mistaken case of worrying about anisotropy in the embedding space in Holtzman et al. (2023).

We do not at all believe that prompting and mechanistic interpretability are at odds with each other, but we do believe that prompting has been denigrated as a form of scientific inquiry to understand what mechanisms LLMs use. We therefore proceed in comparing the goals and methods of what we call “prompt science” and the current state of mechanistic interpretability.




Figure 3: An overview of the comparison between prompt science and interpretability.

5 Prompt Science vs. Mechanistic Interpretability

To compare prompting with mechanistic interpretability, we start with three philosophical perspectives, and then discuss their practical implications.

5.1 Philosophical Perspectives

Different levels of analysis. We start with a perspective inspired by Griffiths et al. (2024). David Marr’s levels of analysis provide a compelling framework for understanding how prompt science and mechanistic interpretability complement each other (Marr, 1982). Marr proposed three levels for understanding complex information-processing systems:

1. *Computational level*: What is the system doing and why?
2. *Algorithmic level*: How does the system do it, in terms of representations and processes?
3. *Implementation level*: How is the system physically realized?

Prompting excels at revealing insights at the computational level—discovering *what* capabilities models have and *why* they might have developed them through training. When we discover that LLMs can perform few-shot learning through careful prompting, we are gaining computational-level insights about the model’s function. When we find what kind of prompts elicit certain behavior, we can often get a sense of what training data these capabilities were influenced by—giving us a peek into the why. For instance, McCoy et al. (2024) show that the model can use shift ciphers with shift values that are commonly used on the internet (1, 3, and 13) but exhibit drastically reduced performance on other values. In comparison, mechanistic interpretability shines at the implementation level, revealing how specific patterns of weights and activations realize particular functions. For instance, mechanistic interpretability researchers used the structure of modulo arithmetic to examine generalization (Liu et al., 2022) and eventually Nanda et al. (2023) fully reverse engineered the algorithm used in small transformers trained on modular addition tasks, which uses discrete Fourier transforms and trigonometric identities to convert addition to rotation about a circle. This led to studies showing that multiple different strategies for modulo arithmetic can be learned (Zhong et al., 2023), and eventually scaled-up studies examining these in medium-sized LLMs (Kantamneni & Tegmark, 2025).

Both approaches meet at the algorithmic level, where we aim to understand the representations and processes models use. Prompting can reveal algorithmic insights through structured behavioral probes, while mechanistic work can in principle identify the specific computational graphs that implement these algorithms. Prompting is better at helping us map out the diversity of possible behavior in many different contexts and is easy to use on models of any scale, where as mechanistic interpretability

shines when attempting to identify small differences (e.g., the clock and the pizza (Zhong et al., 2023)) in processing and remains challenging to scale up.

Different abstraction languages. A fundamental distinction between prompting and mechanistic interpretability lies in their abstraction languages. Mechanistic interpretability relies on geometric abstractions in representation space—circuits, feature directions, and activation patterns. These offer precision but at the cost of having to fully specify definitions, in a way that can limit their applicability. Prompting, by contrast, leverages linguistic abstractions through the manifold of language itself. Importantly, since LLMs are trained on such huge linguistic corpora, we should expect the structure of language to mirror aspects of the model itself.

Language-based abstractions enable what we might call “productive vagueness”—they can specify hypotheses at varying levels of precision, matching our own incomplete understanding of complex phenomena. When we prompt a model to “think step-by-step,” we are not specifying a precisely defined direction in activation space, but rather invoking a rich cluster of associations that the model has learned. This resulting chain-of-thought allows models to expose aspects of what they are doing that would be difficult to describe in model weights.

Crucially, linguistic abstractions naturally align with human conceptual structures, making prompting discoveries more interpretable and actionable for both researchers and practitioners. The tradeoff is precision—linguistic prompts allow a certain ambiguity that mechanistic approaches aim to eliminate. However, there is no sacrifice in scientific rigor, because both abstraction languages—vector representations and prompting—lend themselves to falsifiability with carefully designed experiments. In fact, the ambiguity in language often mirrors the distributed, entangled nature of how concepts are actually represented in neural networks. When we find it difficult to precisely describe a task for an LLM by prompting, that is a powerful signal that the LLM does not have an accurate representation of the necessary concepts to understand the task itself.

Different levels of intervention. Both prompting and mechanistic interpretability can be understood as different approaches to intervention-based science. Mechanistic interpretability intervenes on internal model components—modifying weights, freezing activations, or ablating attention heads—to establish causal relationships between components and behaviors.

Prompting intervenes at the model’s natural interface—the input distribution it was trained on—to similarly establish causal relationships between inputs and behaviors. Both are forms of causal intervention, just at different levels of the system.

This perspective helps explain why prompting has been more immediately productive: it leverages the model’s native interface, which is already optimized to produce meaningful responses to linguistic inputs. The model’s architecture has been specifically designed and trained to transform language inputs into language outputs. Prompting takes advantage of this by using the interface the model was built to be manipulated with.

5.2 Practical Implications

Next, we discuss the practical implications of the above differences along four dimensions: scalability, scope, understanding, and scientific narrative.

Scalability. Mechanistic interpretability struggles to scale with the complexity of modern LLMs. Reverse-engineering billions of parameters is costly if not impossible, and it’s unclear whether the resulting explanations are actionable at scale (Holtzman et al., 2023; Räuker et al., 2023). Prompting, by contrast, is lightweight and deployable at scale, enabling efficient behavioral testing and control without circuit-level access.

Scope. Mechanistic interpretability requires precise hypotheses about model internals, which is often infeasible, especially when the relevant features or circuits are unknown. Moreover, composing multiple edits or identifying the right granularity for interpretation or intervention remains an open challenge. As a result, studies have largely focused on explaining and controlling simple and constrained behavior. Prompting operates at the level of language, giving it broad expressive power to specify desired behaviors or test for failure modes, including subtle or complex ones (Garbacea & Tan, 2025; Rajani et al., 2023).

Understanding. Mechanistic explanations prioritize faithfulness, but their precision can make them opaque—like reading a circuit diagram. The over-emphasis on faithfulness loses track of its pragmatic utility. After all, the only truly “faithful” explanation is the computational graph for a given instance, but we all agree that such a faithful explanation is useless. By contrast, prompting generates behavioral evidence and accessible explanations, serving as a starting point for intuition and hypothesis generation. Prior work has emphasized that explanations are diverse and often shallow (Keil, 2006; Lombrozo, 2006; Tan, 2021; Wilson & Keil, 1998), suggesting a pragmatic role for prompting-based understanding.

Scientific Narrative. Many key discoveries in interpretability begin with behavioral anomalies. For instance, the characterization of induction heads followed observations of copying behavior in Transformers (Al-Rfou et al., 2019; Elhage et al., 2021). Prompting offers a structured hypothesis space to surface such behaviors, guiding mechanistic follow-up. In this light, prompting is not just an evaluation tool; it drives scientific discovery by revealing what is worth interpreting.

In summary, prompting offers both practical advantages and valuable insights comparable to mechanistic interpretability. Rather than competing approaches, they represent complementary methods—analogue to how behavioral psychology and cognitive neuroscience both contribute to understanding human cognition.

6 Addressing Arguments for Dismissing Prompting

There are many arguments for dismissing prompting, we respond to those we think are most frequent and impactful below.

Lack of mathematical formalism. A common critique is that prompting lacks the mathematical rigor expected in machine learning research. The absence of clean equations to describe input-output relationships can make prompting feel less scientifically grounded compared to approaches with elegant mathematical formulations. However, this critique assumes that the most important phenomena in LLMs are best captured through mathematical abstractions we currently have notation and definitions for. As argued in §5.1, the abstractions prompting allows us to use to probe models are more natural for understanding model behavior than those we know how to capture in equations. Much of the point of prompting studies is to discover how we can formalize these: In-Context Learning’s language of input/output demonstrations, Chain-of-Thought’s use of an internal monologue, and DSPy’s method of turning prompt-optimization into programming. The equations we write today—whether describing attention mechanisms or optimization landscapes—often fail to capture the rich behavioral patterns that prompting reveals. Just as plant breeders developed systematic methods for eliciting desired behaviors in crops long before understanding the genetic mechanisms underlying those traits, prompting provides a framework for investigating LLM capabilities that may not yet be amenable to compact mathematical description until we discover “the genetic code” of how models interpret language.

Lack of model training. Machine learning culture is biased towards the belief that meaningful research contributions must involve training new models or developing novel architectures (Birhane et al., 2022; Lipton & Steinhardt, 2019; Sculley et al., 2014, *inter alia*). This perspective can make prompting-based discoveries feel less substantial, as they don’t require the computational resources or optimization-expertise associated with model development. However, this view conflates the prestige of model training with scientific value. Many of the most important insights about LLMs—from in-context learning to chain-of-thought reasoning—were discovered through careful prompting rather than architectural innovation. The field’s reward structure may inadvertently discourage researchers from pursuing behavioral investigations that could yield fundamental insights about these systems, but that’s hardly a good reason to dismiss the scientific contributions prompting has and can give us.

Prompt hacking. Prompting allows for a great deal of variability in how LLMs are deployed (Meincke et al., 2025), creating issues similar to p-hacking (Kosch & Feger, 2025) where one can optimize for the correct prompt to get better scores on an evaluation. Critics argue that this variability should be eliminated to make systems operate like traditional computer programs (Morris, 2024). This represents a category error: it assumes we want to evaluate the LLM “in isolation,” with the prompt getting in the way of that assessment. In reality, LLMs are probability distributions that do

not yield answers to questions; only LLMs combined with prompts and decoding methods constitute complete systems that can be evaluated. Attempting to eliminate this variability through universal prompts misses the point entirely. LLMs naturally require context, and prompting is how we provide that context. **The real issue is insufficient prompt optimization.** If a better prompt can give better results, then that is capability of the system that “runs” on the given prompt, just like how we evaluate program efficiency on specific hardware, not with the program or the hardware alone.

Prompts are brittle. Extensive research has documented “prompt brittleness”—the fact that minor perturbations to prompts can dramatically change LLM outputs (Arora et al., 2022; Sclar et al., 2023; Wallace et al., 2019, *inter alia*). Often this is unintentional, and models fail to do the desired task after small, seemingly meaningless changes. While this certainly complicates model control from an end-user perspective, we argue that this sensitivity actually reflects models’ attempts to infer substantial information from limited context. When LLMs lack sufficient context (from their perspective), they compensate through sensitivity to pragmatic implications, a consequence of training objectives that cause them to model implicit aspects of authorship and persona (Andreas, 2022). Instruction-tuning methods have reduced sensitivity to perturbations (Ouyang et al., 2022) but have also measurably reduced flexibility (Lin et al., 2024). This trade-off highlights that prompt sensitivity is a tool for exploring what factors influence model behavior. **Prompt brittleness may make prompt engineering more difficult, but it makes the science of prompting significantly more potent.** Understanding how human intentions in prompts don’t perfectly transfer to model outputs is a key mechanism for cataloging model capabilities, rather than merely an error to be corrected.

Prompts will not generalize across models. Researchers often worry that prompting discoveries will be too model-specific to provide general insights about LLM behavior. This concern reflects valid scientific instincts about the importance of reproducibility and generalization. However, in practice, many fundamental prompting methods, including in-context learning, retrieval-augmented generation, and chain-of-thought reasoning, have proven robust across model families at large enough scales. Rather than being a weakness, the ability to test generalization across models provides a valuable filter for distinguishing fundamental capabilities from model-specific artifacts.

Prompting requires excessive manual effort. While less explicitly discussed in academic literature, a significant barrier to taking prompting seriously as a research method is the perception that it is fundamentally more labor-intensive and less systematic than other approaches for studying LLMs. This is in contrast to the simultaneous but opposite belief that prompting is “trivial” because it simply involves writing text and observing model outputs. The field has grown accustomed to research methods that can be formulated as optimization problems with clear objective functions, positive examples for fine-tuning, reward signals for reinforcement learning, or compositional operations like task arithmetic (Ilharco et al., 2022). Prompting resists these familiar optimization frameworks. While prompting can technically be viewed as discrete optimization, our very incomplete understanding of the irregular structure of language makes traditional optimization methods useful only in limited cases. Yet this perception of difficulty is often misleading. In practice, prompting frequently enables faster iteration cycles and cheaper exploration compared to methods requiring model retraining, making it an efficient tool for behavioral investigation rather than a cumbersome alternative. This is a difficult statement to prove directly, so we instead offer as evidence all the breakthroughs that prompting has led to, some of which are shown in Figure 2.

7 Future Directions for Making Prompting a Science

In this section, we lay out some concrete future directions for prompting as scientific inquiry.

Prompting as Mechanism Discovery. We believe that prompting will continue to be an effective way to discover LLM capabilities and limitations. Developing a symbiosis between prompting and mechanistic interpretability—where prompting discovers ordered anomalies, mechanistic interpretability formalizes them and suggests new places to look—will be critical for pushing the frontier of LLMs. Example questions include:

- Can we use systematic prompt perturbations to map the precise boundaries where model capabilities break down, revealing the implicit structure of learned knowledge?

- Can we find concepts that only machines, not humans, use (Hewitt et al., 2025; Holtzman et al., 2023)? How can we communicate such concepts to humans?
- Can we use controlled prompt interventions to discover causal dependencies between model capabilities, revealing how different types of knowledge and reasoning interact within LLMs and providing hypotheses for mechanistic interpretability?

Prompting for Effective Control of LLMs. Building on the above insights, prompting can be a productive tool for enabling effective control of LLMs. However, much work is required to make this a science. Example questions include:

- Can we identify a minimal set of “control primitives” in prompt space that can be composed to achieve arbitrary behavioral modifications, analogous to how universal computation emerges from simple logical operations? What are the dimensions of control that are and aren’t possible?
- Most prompting strategies have been constrained within a single long prompt and recent studies show that this is actually more effective than multi-turn prompting (Laban et al., 2025). What is special about long-prompts? What information do multi-turn prompts fail to capture? What makes effective strategies in a multi-turn setting?
- Can we develop adversarially robust prompting strategies that maintain control even when models encounter inputs designed to subvert intended behavior, leveraging multi-model verification (e.g., “is this request harmful?”), reasoning chains, and meta-cognitive prompting?

Prompting as Future Proofing. Whatever future models may look like in architecture, modality, training, or hardware—we will certainly use text to help establish what we want. Focusing on prompting can help us future proof this process. Example questions include:

- What prompting strategies are robust to model updates? What defines a “forward-compatible” prompt? And vice versa, how can we develop models that are backward-compatible with existing prompts?
- Can we use patterns discovered through current prompting research to predict and engineer the behavioral properties of future AI systems, essentially creating a “prompting-informed design principles” framework that shapes how next-generation models should be trained?
- Can we identify universal principles of inter-intelligence communication through prompting that will remain valid even as AI systems become vastly more capable than humans—essentially discovering the fundamental “protocol” of how minds can efficiently communicate across intelligence gaps?

References

- Rami Al-Rfou, Dokook Choe, Noah Constant, Mandy Guo, and Llion Jones. Character-level language modeling with deeper self-attention. In *Proceedings of the AAAI conference on artificial intelligence*, volume 33, pp. 3159–3166, 2019.
- Jacob Andreas. Language models as agent models. In *Findings of the Association for Computational Linguistics: EMNLP 2022*, pp. 5769–5779, 2022.
- Cem Anil, Esin Durmus, Nina Panickssery, Mrinank Sharma, Joe Benton, Sandipan Kundu, Joshua Batson, Meg Tong, Jesse Mu, Daniel Ford, et al. Many-shot jailbreaking. *Advances in Neural Information Processing Systems*, 37:129696–129742, 2024.
- Simran Arora, Avaniika Narayan, Mayee F Chen, Laurel Orr, Neel Guha, Kush Bhatia, Ines Chami, and Christopher Re. Ask me anything: A simple strategy for prompting language models. In *The Eleventh International Conference on Learning Representations*, 2022.
- Yuntao Bai, Saurav Kadavath, Sandipan Kundu, Amanda Askell, Jackson Kernion, Andy Jones, Anna Chen, Anna Goldie, Azalia Mirhoseini, Cameron McKinnon, et al. Constitutional ai: Harmlessness from ai feedback. *arXiv preprint arXiv:2212.08073*, 2022.

- Lukas Berglund, Meg Tong, Max Kaufmann, Mikita Balesni, Asa Cooper Stickland, Tomasz Korbak, and Owain Evans. The reversal curse: LLMs trained on "a is b" fail to learn "b is a". *arXiv preprint arXiv:2309.12288*, 2023.
- Abeba Birhane, Pratyusha Kalluri, Dallas Card, William Agnew, Ravit Dotan, and Michelle Bao. The values encoded in machine learning research. In *Proceedings of the 2022 ACM Conference on Fairness, Accountability, and Transparency*, pp. 173–184, 2022. doi: 10.1145/3531146.3533083.
- Trenton Bricken, Aditya Templeton, Jonathan Batson, Brittany Chen, Adam Jermy, et al. Towards monosemanticity: Decomposing language models with dictionary learning. Transformer Circuits Thread (Anthropic), 2023. Available online at transformer-circuits.pub/2023/monosemantic-features.
- Tom Brown, Benjamin Mann, Nick Ryder, Melanie Subbiah, Jared D Kaplan, Prafulla Dhariwal, Arvind Neelakantan, Pranav Shyam, Girish Sastry, Amanda Askell, et al. Language models are few-shot learners. *Advances in neural information processing systems*, 33:1877–1901, 2020.
- Sébastien Bubeck, Varun Chandrasekaran, Ronen Eldan, Johannes Gehrke, Eric Horvitz, Ece Kamar, Peter Lee, Yin Tat Lee, Yuanzhi Li, Scott Lundberg, et al. Sparks of artificial general intelligence: Early experiments with gpt-4. *arXiv preprint arXiv:2303.12712*, 2023.
- Hoagy Cunningham, Aidan Ewart, Logan Riggs, Robert Huben, and Lee Sharkey. Sparse autoencoders find highly interpretable features in language models. In *Proc. of the Int. Conf. on Learning Representations (ICLR)*, 2024. arXiv:2309.08600 [cs.LG].
- Nelson Elhage, Neel Nanda, Catherine Olsson, Tom Henighan, Nicholas Joseph, Ben Mann, Amanda Askell, Yuntao Bai, Anna Chen, Tom Conerly, et al. A mathematical framework for transformer circuits. *Transformer Circuits Thread*, 1(1):12, 2021.
- Leo Gao, Tom Dupré la Tour, Henk Tillman, Gabriel Goh, Rajan Troll, Alec Radford, Ilya Sutskever, Jan Leike, and Jeffrey Wu. Scaling and evaluating sparse autoencoders. *arXiv preprint arXiv:2406.04093*, 2024.
- Cristina Garbacea and Chenhao Tan. Hyperalign: Interpretable personalized LLM alignment via hypothesis generation, 2025. URL <https://arxiv.org/abs/2505.00038>.
- Thomas L Griffiths, Jian-Qiao Zhu, Erin Grant, and R Thomas McCoy. Bayes in the age of intelligent machines. *Current Directions in Psychological Science*, 33(5):283–291, 2024.
- Daya Guo, Dejian Yang, Haowei Zhang, Junxiao Song, Ruoyu Zhang, Runxin Xu, Qihao Zhu, Shirong Ma, Peiyi Wang, Xiao Bi, et al. Deepseek-r1: Incentivizing reasoning capability in LLMs via reinforcement learning. *arXiv preprint arXiv:2501.12948*, 2025.
- John Hewitt, Robert Geirhos, and Been Kim. We can’t understand AI using our existing vocabulary. *arXiv preprint arXiv:2502.07586*, 2025.
- Ari Holtzman, Peter West, and Luke Zettlemoyer. Generative models as a complex systems science: How can we make sense of large language model behavior? *arXiv preprint arXiv:2308.00189*, 2023.
- Gabriel Ilharco, Marco Tulio Ribeiro, Mitchell Wortsman, Ludwig Schmidt, Hannaneh Hajishirzi, and Ali Farhadi. Editing models with task arithmetic. In *The Eleventh International Conference on Learning Representations*, 2022.
- Eric Jonas and Konrad Paul Kording. Could a neuroscientist understand a microprocessor? *PLoS computational biology*, 13(1):e1005268, 2017.
- Subhash Kantamneni and Max Tegmark. Language models use trigonometry to do addition. *arXiv preprint arXiv:2502.00873*, 2025.
- Frank C Keil. Explanation and understanding. *Annu. Rev. Psychol.*, 57(1):227–254, 2006.

- Omar Khattab, Arnav Singhvi, Paridhi Maheshwari, Zhiyuan Zhang, Keshav Santhanam, Sri Vardhamanan, Saiful Haq, Ashutosh Sharma, Thomas T. Joshi, Hanna Moazam, Heather Miller, Matei Zaharia, and Christopher Potts. Dspy: Compiling declarative language model calls into self-improving pipelines. 2024.
- Takeshi Kojima, Shixiang Shane Gu, Machel Reid, Yutaka Matsuo, and Yusuke Iwasawa. Large language models are zero-shot reasoners. *Advances in neural information processing systems*, 35: 22199–22213, 2022.
- Thomas Kosch and Sebastian Feger. Prompt-hacking: The new p-hacking? *arXiv preprint arXiv:2504.14571*, 2025.
- Philippe Laban, Hiroaki Hayashi, Yingbo Zhou, and Jennifer Neville. Llms get lost in multi-turn conversation. *arXiv preprint arXiv:2505.06120*, 2025.
- Yong Lin, Hangyu Lin, Wei Xiong, Shizhe Diao, Jianmeng Liu, Jipeng Zhang, Rui Pan, Haoxiang Wang, Wenbin Hu, Hanning Zhang, et al. Mitigating the alignment tax of rlhf. In *Proceedings of the 2024 Conference on Empirical Methods in Natural Language Processing*, pp. 580–606, 2024.
- Zachary C. Lipton. The mythos of model interpretability: In machine learning, the concept of interpretability is both important and slippery. *Queue*, 16(3):31–57, 2018. doi: 10.1145/3236386.3241340.
- Zachary Chase Lipton and Jacob Steinhardt. Troubling trends in machine learning scholarship. *Queue*, 17(1):45–77, 2019. doi: 10.1145/3317287.3328534.
- Ziming Liu, Ouail Kitouni, Niklas S Nolte, Eric Michaud, Max Tegmark, and Mike Williams. Towards understanding grokking: An effective theory of representation learning. *Advances in Neural Information Processing Systems*, 35:34651–34663, 2022.
- Tania Lombrozo. The structure and function of explanations. *Trends in cognitive sciences*, 10(10): 464–470, 2006.
- Aleksandar Makelov, Georg Lange, and Neel Nanda. Towards principled evaluations of sparse autoencoders for interpretability and control. In *SeT-LLM Workshop at the International Conference on Learning Representations (ICLR) 2024*, 2024.
- D Marr. The philosophy and the approach. *Visual perception: Essential readings*, pp. 104–123, 1982.
- R Thomas McCoy, Shunyu Yao, Dan Friedman, Mathew D Hardy, and Thomas L Griffiths. Embers of autoregression show how large language models are shaped by the problem they are trained to solve. *Proceedings of the National Academy of Sciences*, 121(41):e2322420121, 2024.
- Lennart Meincke, Ethan Mollick, Lilach Mollick, and Dan Shapiro. Prompting science report 1: Prompt engineering is complicated and contingent. *arXiv preprint arXiv:2503.04818*, 2025.
- Sewon Min, Xinxin Lyu, Ari Holtzman, Mikel Artetxe, Mike Lewis, Hannaneh Hajishirzi, and Luke Zettlemoyer. Rethinking the role of demonstrations: What makes in-context learning work? In *Proceedings of the 2022 Conference on Empirical Methods in Natural Language Processing*, pp. 11048–11064, 2022.
- Meredith Ringel Morris. Prompting considered harmful. *Communications of the ACM*, 67(12):28–30, 2024.
- Neel Nanda, Lawrence Chan, Tom Lieberum, Jess Smith, and Jacob Steinhardt. Progress measures for grokking via mechanistic interpretability, 2023. URL <https://arxiv.org/abs/2301.05217>.
- Harsha Nori, Yin Tat Lee, Sheng Zhang, Dean Carignan, Richard Edgar, Nicolo Fusi, Nicholas King, Jonathan Larson, Yuanzhi Li, Weishung Liu, Renqian Luo, Scott Mayer McKinney, Robert Osazuwa Ness, Hoifung Poon, Tao Qin, Naoto Usuyama, Chris White, and Eric Horvitz. Can generalist foundation models outcompete special-purpose tuning? case study in medicine. *arXiv preprint arXiv:2311.16452*, 2023.
- Chris Olah, Nick Cammarata, Ludwig Schubert, Gabriel Goh, Michael Petrov, and Shan Carter. Zoom in: An introduction to circuits. *Distill*, 5(3):e00024–001, 2020.

- OpenAI. Introducing chatgpt, 11 2022. URL <https://openai.com/index/chatgpt/>. Official release announcement.
- OpenAI. Introducing OpenAI o1. <https://openai.com/index/introducing-openai-o1-preview/>, September 2024. Accessed: 2025-5-21.
- Long Ouyang, Jeffrey Wu, Xu Jiang, Diogo Almeida, Carroll Wainwright, Pamela Mishkin, Chong Zhang, Sandhini Agarwal, Katarina Slama, Alex Ray, et al. Training language models to follow instructions with human feedback. *Advances in neural information processing systems*, 35:27730–27744, 2022.
- Gonçalo Paulo and Nora Belrose. Sparse autoencoders trained on the same data learn different features. *arXiv preprint arXiv:2501.16615*, 2025.
- Jacob Pfau, William Merrill, and Samuel R Bowman. Let’s think dot by dot: Hidden computation in transformer language models. In *First Conference on Language Modeling*, 2024.
- Alec Radford, Jeffrey Wu, Dario Amodei, Daniela Amodei, Jack Clark, Miles Brundage, and Ilya Sutskever. Better language models and their implications. *OpenAI blog*, 1(2), 2019.
- Nazneen Rajani, Nathan Lambert, and Lewis Tunstall. Red-teaming large language models. <https://huggingface.co/blog/red-teaming>, 2023. Accessed: 2025-05-20.
- David Rein, Betty Li Hou, Asa Cooper Stickland, Jackson Petty, Richard Yuanzhe Pang, Julien Dirani, Julian Michael, and Samuel R Bowman. Gpqa: A graduate-level google-proof q&a benchmark. In *First Conference on Language Modeling*.
- Tilman R  uker, Anson Ho, Stephen Casper, and Dylan Hadfield-Menell. Toward transparent ai: A survey on interpreting the inner structures of deep neural networks, 2023. URL <https://arxiv.org/abs/2207.13243>.
- Melanie Sclar, Yejin Choi, Yulia Tsvetkov, and Alane Suhr. Quantifying language models’ sensitivity to spurious features in prompt design or: How i learned to start worrying about prompt formatting. In *The Twelfth International Conference on Learning Representations*, 2023.
- D. Sculley, Gary Holt, Daniel Golovin, Eugene Davydov, Todd Phillips, Dietmar Ebner, Vinay Chaudhary, and Michael Young. Machine learning: The high interest credit card of technical debt. In *SE4ML: Software Engineering for Machine Learning (NIPS 2014 Workshop)*, 2014.
- Lewis Smith, Sen Rajamanoharan, Arthur Conmy, Callum McDougall, Janos Kramar, Tom Lieberum, Rohin Shah, and Neel Nanda. Negative results for sparse autoencoders on downstream tasks and deprioritising sae research. DeepMind Safety Research Blog (Mechanistic Interpretability Team Progress Update), March 2025. URL <https://deepmindsafetyresearch.medium.com/negative-results-for-sparse-autoencoders-on-downstream-tasks-and-deprioritising-sae-research-6cadcf125b9>.
- Chenhao Tan. On the diversity and limits of human explanations. *arXiv preprint arXiv:2106.11988*, 2021.
- Gemini Team, Petko Georgiev, Ving Ian Lei, Ryan Burnell, Libin Bai, Anmol Gulati, Garrett Tanzer, Damien Vincent, Zhufeng Pan, Shibo Wang, et al. Gemini 1.5: Unlocking multimodal understanding across millions of tokens of context. *arXiv preprint arXiv:2403.05530*, 2024.
- Eric Wallace, Shi Feng, Nikhil Kandpal, Matt Gardner, and Sameer Singh. Universal adversarial triggers for attacking and analyzing nlp. In *Proceedings of the 2019 Conference on Empirical Methods in Natural Language Processing and the 9th International Joint Conference on Natural Language Processing (EMNLP-IJCNLP)*, pp. 2153–2162, 2019.
- Yizhong Wang, Yeganeh Kordi, Swaroop Mishra, Alisa Liu, Noah A Smith, Daniel Khashabi, and Hannaneh Hajishirzi. Self-instruct: Aligning language models with self-generated instructions. In *Proceedings of the 61st Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pp. 13484–13508, 2023.

Jason Wei, Xuezhi Wang, Dale Schuurmans, Maarten Bosma, Fei Xia, Ed Chi, Quoc V Le, Denny Zhou, et al. Chain-of-thought prompting elicits reasoning in large language models. *Advances in neural information processing systems*, 35:24824–24837, 2022.

Robert A Wilson and Frank Keil. The shadows and shallows of explanation. *Minds and machines*, 8: 137–159, 1998.

Shunyu Yao, Dian Yu, Jeffrey Zhao, Izhak Shafran, Tom Griffiths, Yuan Cao, and Karthik Narasimhan. Tree of thoughts: Deliberate problem solving with large language models. *Advances in neural information processing systems*, 36:11809–11822, 2023.

Tal Yarkoni. The generalizability crisis. *Behavioral and Brain Sciences*, 45:e1, 2022.

Ziqian Zhong, Ziming Liu, Max Tegmark, and Jacob Andreas. The clock and the pizza: Two stories in mechanistic explanation of neural networks. *Advances in neural information processing systems*, 36:27223–27250, 2023.